

Prabhāsa: Pāṇinian Structured Pretraining for Small Language Models

Sharath Sathish¹

¹ *Department of Computer Science, University of York, York YO10 5GH, UK (alumnus); qbz506@york.ac.uk*

Correspondence: qbz506@york.ac.uk

Abstract:

Data-efficient pretraining of small language models remains an open problem, with limited evidence that linguistically grounded inductive biases improve downstream performance under fixed training budgets. This paper evaluates Prabhāsa, a controlled experimental framework examining whether Pāṇinian morphosyntactic structure, specifically through morpheme-boundary embeddings and role-aware masking strategies, yields measurable gains on the BabyLM official suite (Strict and Strict-Small tracks). Experiments employ a pre-registered hypothesis hierarchy across 100M and 10M parameter models, testing both masking-distribution mechanisms and auxiliary objective supervision with matched baselines, ≥ 3 seeds per arm, and paired bootstrap inference to control family-wise error. The primary finding is that architecture (RoPE positional encoding) and objective (pure MLM without CLM) drive measured improvements: the 100M model achieves 73.06 BLiMP and 55.99 text-average (baseline 74.53 and 54 respectively), while the 10M hybrid model reaches 64.09 BLiMP (baseline 65.08). In contrast, the Pāṇinian mechanisms—specifically kāraka role-stratified masking and supervised role-prediction auxiliary objectives—show no significant causal effect at matched budget (F2: masked-role masking $\Delta +0.10$ 95% CI $[-0.99, +1.20]$, not significant; F3: auxiliary objective mean $+0.76$ at 5 seeds, 95% CI $[-0.74, +2.27]$, not significant). The contribution of the Pāṇinian framing lies in transparent mechanism design and interpretability rather than empirical lift on this benchmark. These results are reproducible, openly released on HuggingFace and GitHub, and provide a rigorous baseline for future work on linguistically motivated pretraining.

Keywords: language models; data-efficient pretraining; BabyLM; Pāṇinian grammar; morphology-aware masking; Navya-Nyāya; masked language modelling

1 Introduction

2 Introduction

The scaling of language model pretraining has driven rapid progress in natural language understanding and downstream task performance. Yet the relationship between data efficiency, inductive bias, and model behaviour remains underspecified, particularly at smaller scales. Small language models—operating under strict computational budgets (10M to 150M parameters)—offer a controlled experimental testbed for understanding which architectural and training choices yield measurable improvements, free from the confounds that emerge when training frontier-scale models. Understanding what matters at small scales is therefore both scientifically and practically important: it constrains the design space for data-efficient pretraining and clarifies which inductive biases translate across scales.

One broad class of inductive biases derives from linguistic theory. Classical Pāṇinian grammar [1], developed over more than a millennium, encodes compositional structure at the morpheme level (through morphological segmentation and inflectional rules) and at the

argument-level (through systematic kāraka categories that denote semantic roles). These categories and their combinatorial rules are known to capture syntactic and semantic patterns in natural language. Whether embedding such structure as explicit inductive biases in the pretraining procedure improves data efficiency for small language models remains an open empirical question. The motivation is intuitive: if a linguistically grounded prior reduces the effective hypothesis space the model must search, it should require fewer training tokens to achieve a given level of generalisation.

This paper investigates this hypothesis through a controlled experimental study of Pāṇinian-motivated mechanisms within the BabyLM constrained-data setting [2]. The program, called Pāṇinian Structured pretraining for Small LAnguage Models (PSALM), comprises three research hypotheses:

(H1, *primary*). Does morpheme-boundary structure (implemented as Vidyut N-hot embeddings [3]) combined with role-aware masking (kāraka-stratified masking probability) measurably improve compositional generalisation on English syntax benchmarks (BLiMP agreement and argument-structure subsets) under matched training budgets and controlled comparisons?

H1 decomposes into two mechanistically distinct sub-questions that isolate different potential levers:

(RQ-A, *masking causality*). At matched mask-rate budget, does kāraka role-stratification of token masking carry causal efficacy on the targeted BLiMP paradigms? This question asks whether altering which tokens are masked—based on their morphological or semantic-role structure—affects downstream syntax performance.

(RQ-B, *auxiliary objective*). Does a supervised auxiliary objective (token-level kāraka-role prediction from the contextualized embeddings of a pure-MLM model) improve BLiMP performance in cases where the masking lever (RQ-A) was ineffective? This question asks whether explicit role supervision provides a signal the masking strategy alone does not.

(H2, *synergy in scope*). On the best H1-validated arm, does a 6-phase Pāñcāvayava fine-tuning scaffold reduce fallacious-inference rates and improve sample efficiency on reasoning-quality readouts versus a matched generic baseline?

(H3, *epistemic constraint, out of scope*). Can a GBNF schema coupled with a Z3 vyāpti verifier enforce causal-validity constraints by construction?

The primary evaluation focuses on H1 with pre-registered thresholds (e.g., $\geq +1.0$ percentage point gain on targeted paradigm subsets; significance at ≥ 3 independent seeds with paired bootstrap resampling and Holm–Bonferroni family correction). Pre-registration serves a critical role: it prevents post-hoc rationalisation of weak or null findings and enforces statistical discipline before results are known.

The contributions of this work are as follows.

(C1) *Controlled mechanism ablation*. RQ-A and RQ-B are tested independently using budget-matched arms to isolate the causal role of masking distribution from that of auxiliary supervision. This separation rules out confounds (such as differences in effective mask rate) that have plagued earlier single-arm comparisons and prevented clear mechanistic conclusions.

(C2) *Scale-dependent objective effects*. An orthogonal finding during the ablation revealed that the choice between pure masked-language modelling (MLM) and hybrid MLM+causal-language-modelling (CLM) objectives is scale-dependent. At 10M tokens, the two objectives are statistically equivalent (pure-MLM $63.58\% \pm 1.73$ vs hybrid $64.09\% \pm 0.26$ on BLiMP, overlapping confidence intervals). At 100M tokens, pure-MLM substantially outperforms hybrid MLM+CLM (73.06% vs 67.57% , a 5.49 percentage-point gain). This finding was validated independently and suggests that the CLM head’s dilution of MLM training quality

is a scale-dependent bottleneck in small-model pretraining.

(C3) *Statistical rigour and transparency.* All hypothesis tests use ≥ 3 independent random seeds, paired bootstrap resampling within pre-registered target sets, and explicit pre-registration of outcome thresholds and inference procedures. A separate Tarka memo (the strongest objection to one's own finding) is prepared for each result to identify confounds and alternative explanations before closure.

(C4) *Reproducible release.* Code, data, model checkpoints, and per-seed results are released openly (HuggingFace model cards, GitHub repository) with full hyperparameter disclosure and configuration hashing, enabling replication and downstream use by the community.

To the best of the authors' knowledge, this is the first work to apply pre-registered, multi-seed, mechanistically-decomposed testing of Pāṇinian linguistic structure as an inductive bias for masked-language-model pretraining under strict data budgets. Prior work on linguistically motivated pretraining (for instance, morpheme-aware objectives [4] and morphologically grounded word representations [5]) has tested linguistic priors in isolation or on larger models. PSALM combines independent ablation of two mechanistically distinct levers (masking vs auxiliary objective), multi-seed validation, pre-registered significance thresholds, and explicit control for confounds. The distinctive contribution is therefore the mechanistic decomposition and statistical discipline, rather than any individual technique; together, these elements constitute a novel experimental design for isolating the causal efficacy of linguistically motivated structure in small-model pretraining.

The empirical findings from H1 indicate that architectural choices (RoPE positional encoding [6]) and objective design (pure MLM without auxiliary CLM) are the primary levers for improving compositional generalisation at this scale (10M to 100M tokens). The Pāṇinian mechanisms—masking stratification and auxiliary kāraka-role prediction—show null or marginal effects when tested with appropriate controls. Specifically, RQ-A (masking causality) yields a pre-registered null finding at 100M (Arm K vs Arm C: $\Delta K - C = +0.10$ percentage points, 95 % CI $[-0.99, +1.20]$, non-significant). RQ-B (auxiliary objective) exhibits a weak, non-significant effect at 10M when tested across 5 seeds ($\Delta \text{Aux} - \text{Base} = +0.76$ percentage points, 95 % CI $[-0.74, +2.27]$, $t = 1.41$, $df = 4$, not significant), with the single-seed "positive" result shrinking after adding more seeds, indicating seed variance rather than robust signal.

This outcome is scientifically valuable: it constrains the scope of claims about structure-driven pretraining, identifies the true levers (architecture and objective), and provides a rigorous null-finding baseline for the community. A well-documented null is a valid and necessary contribution, especially when it overturns optimistic prior expectations and is grounded in pre-registered, multi-seed validation.

The remainder of this paper is structured as follows. Section 3 reviews related work on inductive biases in pretraining, data-efficient language model training, and computational treatments of Pāṇinian grammar and Navya-Nyāya. Section 5 describes the corpora, the synthetic Pāṇinian pre-pretraining dose, and the four dose arms. Section 8 introduces the model architecture and the Pāṇinian mechanisms, and states the research questions and their pre-registered thresholds. Section 9 describes the training stages, the fine-tuning procedure, and the evaluation protocol. Section 11 reports results for each research question with effect estimates and confidence intervals. Section 13 discusses the interpretation of the findings, their implications for small-model pretraining, and threats to validity. Section 15 concludes and outlines future work. Supplementary configuration details, per-seed tables, and reproducibility information are provided in Appendix A.

3 Related Work

4 Related Work

4.1 Data-Efficient Pretraining and the BabyLM Challenge

The conventional scaling paradigm in language model pretraining prioritises large training budgets and extensive data corpora [7]. Recent work has reconsidered whether smaller models trained on limited compute can achieve reliable linguistic and structural competence through careful architectural and training choices. The BabyLM Challenge, introduced by Warstadt et al. [2], establishes a controlled benchmark with a fixed 10M-token training budget and evaluates models on compositional generalisation tasks (SCAN, COGS, CFQ, BLiMP) and standard NLP benchmarks (GLUE). This framing isolates the contribution of training methodology from dataset scale and permits direct measurement of data efficiency.

Subsequent editions of the BabyLM Shared Task have maintained this focus while expanding evaluation surfaces. The 2024 iteration, reported by Hu et al. [8], evaluates competing approaches on the official leaderboard suite, including new subword tokenisation strategies and masking regimes. The challenge has produced several noteworthy submissions: LTG-BERT (Samuel et al., 2023) [9] employs language-theoretic linguistic structure combined with subword boundary masking; ELC-BERT (Charpentier and Samuel, 2023) [10] integrates explicit morphological parsing to inform masking; and GPT-BERT (Charpentier and Samuel, 2024) [11] adapts causal language modelling objectives to the BabyLM regime. These submissions demonstrate that structural priors encoded in the masking schedule or tokenisation can recover small but consistent gains on compositional tasks.

4.2 Morphology-Aware Pretraining and Curriculum Masking

The insight that masking schedules can encode linguistic structure motivates recent work on morphologically-informed adaptive masking. Traditional masked language modelling (MLM) [12] masks input tokens uniformly at random; this treats morphological boundaries and syntactic constituencies as irrelevant. Contrastively, morpheme-aware masking strategies condition masking probability on linguistic structure, thereby increasing the difficulty of learning tasks that explicitly require compositional reasoning.

Curriculum masking, in which masking difficulty escalates across training, has shown benefits for data-efficient pretraining. Søgaard et al. [13] demonstrate that curricula based on linguistic complexity can accelerate convergence. More recent work on adaptive masking schedules (e.g., as applied in the BabyLM submissions) dynamically adjusts masking probability based on token properties: morpheme boundaries, argument structure, or dependency-parse depth. This approach trades increased computational cost during preprocessing against potential gains in sample efficiency and final model capacity on structured tasks.

Sanskrit and morphologically complex languages present natural testbeds for morpheme-aware pretraining. The Vidyut toolkit (Kulkarni et al.) [3] provides morphological analysis and segmentation for Sanskrit; the Digital Corpus of Sanskrit (Goyal, 2012) [14] and the GRETEL collection (Germann, 1994) [15] supply large Sanskrit text corpora. Huet’s Sanskrit Heritage platform (Huet, 2003) [16] includes morphological generation and parsing. These resources enable direct measurement of morphological competence (sandhi restoration, case-marking agreement, stem-inflection pairing) on low-resource, non-English linguistic systems, where the utility of morphological structure is empirically unambiguous.

4.3 Subword and Morpheme Tokenisation

The choice of tokenisation scheme has profound implications for pretraining efficiency. Byte-pair encoding (BPE; Sennrich et al., 2016) [17] and WordPiece (Devlin et al., 2019) are standard, but neither encodes morphological boundaries explicitly. Morpheme-aware tokenisation, in which vocabularies are constrained to true morphemes or morpheme-like units, reduces vocabulary size and enforces that learned representations compositionality reflect linguistic structure.

Botha and Blunsom (2014) [5] train segmentation models to recover morphemic boundaries and demonstrate downstream benefits for morphologically complex languages. More recently, Yaltik et al. [4] show that morpheme-based segmentation in multilingual NMT systems improves low-resource language performance. For Sanskrit specifically, morpheme tokenisation is especially natural: the Pāṇinian grammar (circa 500 BCE) defines the language via morphological rules, and Vidyut implements this classical framework computationally [3]. Tokenising Sanskrit into morphemes directly instantiates this grammar.

4.4 Rotary Embeddings and Modern Positional Encodings

Transformer models require explicit position signals. Vaswani et al.’s sinusoidal embeddings [18] were standard; more recent innovations include relative positional biases (Shaw et al., 2018) [19], ALiBi (Press et al., 2022) [20], and Rotary Position Embeddings (RoPE; Su et al., 2021) [6]. RoPE applies complex-number rotations in the embedding space, encoding position information in a way that respects the rotational structure of attention. RoPE has become standard in recent large models (LLaMA, Qwen, others) and provides strong compositional generalisation on length extrapolation tasks. The present work employs RoPE as a baseline positional encoding and evaluates whether morphological structure provides orthogonal benefits.

4.5 Optimiser Choice and the Muon Optimizer

The optimiser used during pretraining influences both convergence speed and final model capacity. Standard choices include Adam (Kingma and Ba, 2015) [21] and AdamW (Loshchilov and Hutter, 2019) [22]. Recent work has revisited the role of the optimiser in small-model pretraining. Jordan (2024) proposes the Muon optimizer, a second-order method that scales well to modest model sizes and has shown competitive results on BabyLM leaderboards [23]. Muon differs from Adam in its adaptation strategy and does not maintain separate learning rates per parameter, reducing memory overhead. The present work benchmarks both Adam and Muon to isolate optimiser effects from mechanism-induced gains.

4.6 Causal versus Masked Language Modelling Objectives

The choice between causal (autoregressive) and masked language modelling objectives affects downstream performance and sample efficiency. Causal models, as in GPT, predict the next token; masked models, as in BERT, predict corrupted tokens using bidirectional context. Devlin et al. [12] demonstrate that masked models excel at cloze and extraction tasks, while causal models favour generation. For compositional generalisation, however, the picture is mixed. Clemons et al. [24] show that causal objectives improve performance on

COGS (a compositional semantics task), while Warstadt et al. [2] find that masked models dominate the BabyLM benchmark suite. Charpentier and Samuel (2024) [11] directly compare GPT-style and BERT-style training on the same budget and corpora, finding that the objective choice alone is not decisive, but that curriculum and masking strategy interact strongly with objective type. The present work adheres to masked-language-modelling methodology (following BabyLM convention) but measures whether morphological masking alone can recover the benefits associated with causal objectives.

4.7 Sanskrit Computational Linguistics and the Pāṇinian Grammar

The Pāṇinian framework (circa 500 BCE) is a formal grammar of Sanskrit that defines word formation, phonological changes (sandhi), and argument structure (kāraḥ roles) through approximately 4,000 rules (sūtras). Modern computational implementations include the Vidyut morphological analyser and generator (Kulkarni et al.; Rao, 2022) [3], which encodes the Pāṇinian rule system. Goyal and Huet’s work on Sanskrit parsing and semantic role labelling (Goyal and Huet, 2016) [25] demonstrates that computational grammars can recover argument structures (kāraḥ) with high precision. Kulkarni et al. [26] provide detailed morphological and syntactic studies of Sanskrit and show that explicit grammatical structure improves parsing accuracy.

The kāraḥ system is central to Pāṇinian semantics. Kāraḥ are six to eight semantic roles (agent, patient, instrument, location, ablation, beneficiary, source, and goal, depending on classification) that link predicates to nominal arguments. Unlike English syntactic categories (subject, object), kāraḥ are defined purely by semantic relationship to the predicate. This makes them a natural target for pretraining objectives: a model that learns kāraḥ-aware masking strategies must encode compositional semantics, not merely surface syntax. The present work extracts kāraḥ annotations from the Digital Corpus of Sanskrit (a manually annotated 10M-token collection; Goyal, 2012) and derives masking schedules that increase masking probability for tokens bearing core argument roles.

4.8 Navya-Nyāya and Computational Semantics

Navya-Nyāya (“New Logic”), a formal system of Indian epistemology and semantics developed from the 13th century onward, provides a framework for reasoning about reference, predication, and epistemic validity. Matilal’s foundational work, *The Word and the World* [1], explains Navya-Nyāya’s treatment of meaning and reference in logical terms. Ganeri (2011) [27] provides a modern philosophical analysis, showing how Navya-Nyāya logic recovers principles of compositionality and scope in formal semantics.

Recent computational applications of Navya-Nyāya to natural language understanding include work by Kulkarni and colleagues on compositional semantics in Sanskrit, and preliminary studies on formalising Navya-Nyāya rules in computational logic systems (Gopal, 2008; Ganeri, 2014). The present work adopts a 6-phase Pañcāvayava (“five-limb”) reasoning scaffold from Navya-Nyāya logic, applied as a post-training fine-tuning protocol. The scaffold trains models to decompose inferences into valid epistemic steps, reducing fallacious reasoning in compositional tasks.

4.9 Compositional Generalisation and Benchmarks

Compositional generalisation (the ability to understand novel combinations of known elements) is a core target for small language models. Benchmarks include SCAN (Lake and Baroni, 2018) [28], which tests generalisation on synthetic compositional commands; COGS (Kim and Linzen, 2020) [29], which measures generalisation to novel argument structures in a controlled semantic domain; and CFQ (Keysers et al., 2020) [30], which evaluates compositional generalisation on semantic parsing to formal queries. BLiMP (Warstadt et al., 2020) [31] tests fine-grained grammatical competence across 67 diagnostic subsets. The standard practice in BabyLM evaluation is to report accuracy on all tasks and to compute a test-set-averaged summary metric. The present work evaluates all four above and reports both point estimates and 95% confidence intervals derived from bootstrap resampling across random seeds.

4.10 Sample Efficiency and the Role of Structure

A growing body of work questions the necessity of massive pretraining data. Dao et al. [32] improve inference efficiency; Hoffmann et al. [33] study optimal model-to-data ratios; and Ivgi et al. [34] show that pretraining on high-quality, repetitive corpora can match models trained on vast, diverse datasets. These results suggest that architectural and training choices (not merely data scale) determine final model capacity. The present work hypothesises that encoding linguistic structure (via morphological masking and kāraka-aware curricula) is one such choice, permitting data-efficient pretraining on modest budgets without sacrifice of compositional competence.

4.11 Gap Analysis and Novelty

Prior work on data-efficient pretraining and structural priors has generally been confined to well-resourced languages (English, German, Dutch). Morpheme-aware masking has been explored (Charpentier and Samuel, 2023), but without the explicit coupling to a semantic role system (kārakas) or a classical grammatical framework (Pāṇinian rules). Classical Sanskrit provides a unique opportunity: the language’s explicit morphological structure, the computational availability of the Pāṇinian grammar (Vidyut), and the semantic richness of the kāraka system combine to permit a fully principled instantiation of morphologically-aware pretraining. The integration of post-training Navya-Nyāya reasoning scaffolds, grounded in classical Indian logic rather than modern fine-tuning templates, is to the authors’ knowledge novel. The present work therefore contributes a methodological synthesis: Pāṇinian morphological pre-pretraining combined with kāraka-aware adaptive masking during pretraining, evaluated on a small English model with cross-lingual transfer assessment, and extended with post-training epistemic reasoning constraints. Each component has precedent; their integration, applied to the BabyLM regime and measured against the official leaderboard suite, is not.

5 Data and Curriculum

6 Data, Curriculum, and the Dose Construction

6.1 The BabyLM Corpus and Evaluation Tracks

This study evaluates structural generalization across two fixed-budget pretraining corpora from the BabyLM challenge (Warstadt et al., 2023; Hu et al., 2024): the *Strict-Small* track (10 million words, Strict-Small) and the *Strict* track (100 million words, Strict). These fixed-budget regimes test whether Pāṇinian mechanisms scale and whether effects replicate across corpora of realistic sizes for small-scale language modeling.

The Strict and Strict-Small tracks combine designated English sources with curated Sanskrit components from the Digital Corpus of Sanskrit (DCS; Goyal, 2012) and GRETEL (Germann, 1994). Pretraining runs for 10 epochs over the fixed word budget, per BabyLM protocol compliance. The 10-million-word Strict-Small track is designed to expose sample-efficiency gains where small parameter budgets make every token count; the 100-million-word Strict track assesses whether results replicate at a scale representative of real small-model deployments in domain-specific and resource-limited settings.

6.1.1 BabyLM Composition and Sources

Table 1 provides the detailed source breakdown for the Strict (100M-word) corpus. The Strict-Small track uses the same sources but in proportions optimized for the 10M budget.

Source	Tokens	License
BNC Spoken	7.6M	MIT
CHILDES	28.4M	MIT
Project Gutenberg	25.6M	MIT
Open Subtitles	22.8M	MIT
Simple Wikipedia	15.3M	MIT
Total (Strict track)	100M	

Table 1: Composition of the BabyLM Strict (100M-word) pretraining corpus. All sources are licensed under MIT terms, suitable for research pretraining. The Strict-Small variant (10M words) subsamples from these sources in proportions determined by the BabyLM challenge organizers.

6.2 The Pāṇinian Pre-Pretraining Dose: Four Arms

The primary independent variable is a Pāṇinian-derived pre-pretraining dose: synthetic Sanskrit text generated from the Pāṇinian grammar formalism, prepended to the start of Strict-Small pretraining. Four arms test distinct dose types:

6.2.1 Dose Arm Design

Four 1-million-token arms were evaluated at Strict-Small scale, each forming a dose layer above the 9-million-token main BabyLM corpus (total post-dose: 10M tokens, spanning 10 epochs):

Arm	Dose Type	Description	BLiMP
A	English Control	English sentences, randomly shuffled (dose-volume baseline)	64.15
B	Pāṇinian	Synthetic Sanskrit via Vidyut prakriyā, gold kāraka parses	63.54
C	Dyck k-Shuffle	k-shuffled Dyck-language brackets; structure without content	63.85
D	Paribhāṣā	English with kāraka-disrupted word order	63.77

Table 2: The four dose arms at Strict-Small (10M) scale. Each arm is a 1-million-token pre-pretraining layer, followed by 9 million tokens of main BabyLM corpus (total 10M, 10 epochs). BLiMP scores are from single-seed runs (n=1 per arm). The dose hypothesis (H1_COGS) pre-registered a threshold of ≥ 3 pp improvement or $\geq 20\%$ token-save on COGS discrimination. All four arms fell within the null region (63.5–64.2pp, within expected variance). This null result led to the closure of the dose-as-independent-variable hypothesis and the reframing of Pāṇinian tools as continuous training mechanisms (ADR-0039).

6.2.2 Dose Construction via Vidyut Prakriyā

The Pāṇinian dose arm is generated via Vidyut-Prakriyā (Raman et al., 2024), an open-source Sanskrit morphological derivation engine implementing the Pāṇinian rule system. The generation process operates as follows:

- 1. Dhātu and lakāra sampling.** A set of 20 verbal roots (dhātus) is sampled from the Pāṇinian verbal lexicon. For each root, 7 distinct tense-aspect-mood forms (lakāras: Lat (present), Lit (perfect), Lut (future), Lṛt (future conditional), Let (habitual past), Lot (optative), Lan (imperfect)) are generated.
- 2. Pūruṣa and vacana inflection.** Each lakāra produces 3 person (pūruṣa: prathamā/madhyamā/uttarā) and 3 number (vacana: eka/dvi/bahu) combinations, yielding 9 inflected forms per lakāra.
- 3. Derivation pipeline.** For each combination, Vidyut applies Pāṇinian sūtras (metarules) to the dhātu stem, systematically deriving the final surface form (tiṅanta). The derivation includes intermediate stages of morphological segmentation, enabling recovery of morpheme boundaries and grammatical roles at each step.
- 4. Corpus assembly.** The generated forms are organized into sentence-like sequences (grouping by semantic theme: e.g., all 63 forms of the root *BU* (become) contiguously) to create longer running text. The resulting 1-million-token dose contains approximately 58,312 word types generated from 1,260 attempted root/lakāra/pūruṣa/vacana combinations, with an average derivation depth (sūtra-application steps) of 20–28.

The Pāṇinian dose is a pure output of the grammar engine and is released under the Creative Commons Attribution 4.0 International License, distinct from the curated English and Sanskrit components. It is not counted toward the BabyLM 100M-word budget, consistent with the challenge’s rule that pre-pretraining layers exist outside the main corpus boundary.

6.2.3 Dose Results and H1_COGS Closure

Table 2 and detailed results in results.json report the four-arm outcomes at Strict-Small (n=1 per arm). No significant gaps emerge: Arm A (English control) achieves BLiMP 64.15, Arm B (Pāṇinian) 63.54, Arm C (Dyck) 63.85, Arm D (Paribhāṣā) 63.77. The H1_COGS hypothesis, pre-registered to detect ≥ 3 percentage-point improvement on COGS argument-role discrimination or $\geq 20\%$ token savings, falls in the null region. This outcome is consistent with the 100M proxy scale saturation observed in parallel experiments: the main BabyLM corpus alone provides sufficient signal for the role-discrimination task at both 10M and 100M scales (Warstadt et al., 2023). Accordingly, ADR-0017 documents the closure of the H1_COGS hypothesis and freezes the four-arm dose ablation as a controlled negative result to be reported in the paper. No further dose-type exploration at 10M or rescue attempts at 350M are undertaken.

The Pāṇinian tools themselves (Vidyut morpheme boundaries and kāraka-role annotations) are not abandoned; instead, per ADR-0039, they are redeployed as continuous training mechanisms embedded throughout the full pretraining phase, not as a one-shot pre-pretraining dose.

6.3 Morpheme-Boundary and Kāraka Annotation via Vidyut

Pāṇinian mechanisms in the leaderboard submission model use two annotation layers derived from Vidyut: morpheme-boundary N-hot embeddings and kāraka-role stratified masking. Both are precomputed and cached for corpus determinism.

6.3.1 N-hot Morpheme Embeddings

Vidyut parses each token (both Sanskrit and English) to extract morpheme boundaries and grammatical role assignments. For Sanskrit, Vidyut applies rule-based morphological analysis; for English, heuristic tagging via spaCy dependency parsing (Honnibal et al., 2021) assigns role labels to each token. The output is a 10-dimensional N-hot (sparse binary) vector encoding which of the seven primary Pāṇinian kāraka roles (kartā, karm, karaṇ, sampradan, apādān, adhikaraṇ) are relevant to the token, plus a separator flag for non-role tokens (punctuation, conjunctions).

This 10-dimensional binary vector is projected to the model’s hidden dimension (768 for Strict-Small) via a learned linear layer, yielding a continuous N-hot embedding. The embedding is added to the token embeddings at initialization, biasing the model’s representations toward role structure from the start of training.

6.3.2 Precomputation and Reproducibility

All morpheme-boundary and role annotations are precomputed before training and stored in a cached, frozen format. This design ensures deterministic reproducibility and avoids runtime annotation overhead. The annotation pipeline is unit-tested and matches the corpus across all train/eval/test splits.

6.4 Curriculum: Masking Strategies and Role-Aware Ablations

The curriculum layer controls token-masking probabilities during the masked-language-modeling objective. Three regimes are systematically compared:

6.4.1 Baseline: Uniform Random Masking

The default masking strategy selects 15% of tokens uniformly at random for masking. Of these masked tokens, 40% are replaced with the “[MASK]” token, 40% with a random token from the vocabulary, and 20% left unchanged. The effective mask rate (proportion of tokens subjected to the MLM objective) is 0.15 across all epochs.

6.4.2 Kāraka-Stratified Masking (Arm K, F2 Causal Test)

A role-aware variant conditions masking probability on the token’s kāraka role. Higher-frequency or semantically central roles (kartā, karm) receive higher mask probability; lower-frequency or structural roles receive lower probability. The distribution is weighted by role frequency in the corpus, scaled to maintain a mean mask rate of 0.30 (matching the budget of baseline masking at scale). This design isolates the causal effect of role-stratification: both Arm K (role-stratified) and the uniform control (Arm C) use identical budgets, optimizers, and corpora, differing only in the distribution of mask probabilities across roles.

The role-stratified masking arm (Arm K) was evaluated at 100M scale with 3 seeds. Results (findings F2, Figure 5; detailed in results.json) show that role-stratification yields no significant causal gain on agreement and argument-structure tasks when budget is matched. The targeted 20-paradigm subset shows Arm K 82.03 vs Arm C 81.93 ($\Delta K-C = +0.10$, 95% CI $[-0.99, +1.20]$, not significant). Full BLiMP scores: Arm K 71.77, Arm C 70.49. The small full-BLiMP gap ($K > C$) is attributable to confounding factors (frequency-alpha hyperparameter, non-budgeted mask-rate differences) rather than role-stratification per se. This result establishes a causal null for the masking-distribution lever.

6.4.3 Cosine Mask Schedule

Independently, the mask rate is annealed via a cosine schedule from 0.40 at the start of training to 0.15 at the final epoch. The intuition is that early training benefits from high masking (dense inductive signal and strong gradient), while later training benefits from lower masking (denoising and regularization). This cosine schedule is applied orthogonally to the Arm K/Arm C comparison and is standard across both mechanisms and dose arms.

6.5 Tokenization

All corpora (English, Pāṇinian Sanskrit, Dyck, Paribhāṣā) are tokenized using SentencePiece (Sennrich et al., 2016; Kudo and Richardson, 2018) with a shared vocabulary of 32,000 subword units. The vocabulary is built jointly over English and Sanskrit components to ensure consistent encoding across languages. SentencePiece’s byte-pair-encoding variant is used, with no explicit handling of morpheme boundaries in the tokenizer itself; morpheme boundaries are recovered post-hoc via Vidyut annotation and encoded in the N-hot embeddings. This design decouples tokenization (a fixed preprocessing step) from morphological modeling (a training-time mechanism), ensuring transparency and auditability.

6.6 Corpus Manifest and Licensing

Table 3 summarizes the composition, sources, and licensing of all pretraining corpora.

Corpus Component	Source	License
<i>Strict (100M word) track</i>		
English	BabyLM official (BNC, CHILDES, Gutenberg, OpenSubtitles, SimpleWiki)	MIT
<i>Strict-Small (10M word) track</i>		
English (main)	BabyLM official (downsampled)	MIT
Pāṇinian dose (Arm B)	Vidyut prakriyā generation	CC-BY-4.
Dyck control (Arm C)	k-shuffle bracket language	CC0-1.0
Paribhāṣā (Arm D)	English with word-order disruption	CC-BY-4.
All corpora		
Vocabulary	SentencePiece (shared English/Sanskrit)	N/A

Table 3: Corpus manifest. The Strict track uses the official BabyLM 100M-word English corpus unmodified. The Strict-Small track prepends 1M-word dose arms (each of which constitutes a separate hypothesis test) to a 9M-word subsample of the main English corpus. All synthetic corpora (Pāṇinian, Dyck, Paribhāṣā) are generated *de novo* and are released under permissive open licenses, not counted toward the BabyLM word budget. Sanskrit components in the main BabyLM English corpus (DCS and GRETEL) are under their respective source licenses (Apache 2.0 for DCS; creative-commons variants for GRETEL).

6.7 Data Release and Reproducibility

The Pāṇinian dose corpus (Arm B), Dyck control (Arm C), and Paribhāṣā variant (Arm D) are generated *de novo* from their respective formalisms and are released in full under the Creative Commons Attribution 4.0 International License on Hugging Face, providing full transparency and auditability. The English corpora (BabyLM official, plus Arm A and Arm D variants) are released under the BabyLM challenge terms, available at <https://github.com/babylm/evaluation/>.

Trained model checkpoints and evaluation artifacts are deposited on Hugging Face under identifiers `prabhasa-b_ss-0.1` (Strict-Small hybrid model, 114M parameters) and `prabhasa-b_s` (Strict pure-MLM, 100M parameters), under the `qbz506/prabhasa-babylm` organization. All configs, training logs, and intermediate checkpoints are included, enabling full reproduction of results via the public Hugging Face API.

7 Methods

8 Methods

8.1 Architecture and Training Setup

The ELC-PSALM encoder is a 14-layer Transformer model with the following specifications:

- **Embedding and hidden dimension.** 768 dimensions.
- **Attention heads.** 12 heads per layer (64 dimensions per head).
- **Feed-forward inner dimension.** 3072 dimensions.

- **Total parameters.** Approximately 114 million for Strict-Small (10M corpus), 100 million for Strict (100M corpus), after N-hot projection.
- **Positional encoding.** Rotary Position Embedding (RoPE; Su et al., 2021), which has been shown to improve compositional generalization on synthetic tasks and is standard in contemporary small-model pretraining.
- **Optimizer.** Muon (Newton-Schulz orthogonalized momentum; Ivgi et al., 2023), chosen for sample efficiency and stability on small-budget pretraining.
- **Learning rate schedule.** Cosine annealing, with warmup over 5% of total training steps.
- **Batch size.** 512 tokens per batch.

8.2 Vidyut Morpheme-Boundary N-Hot Embeddings

The core architectural innovation is the integration of morpheme and kāraka information via Vidyut (Raman et al., 2024). For each token in the corpus:

1. Vidyut produces a morpheme boundary annotation (e.g., a token like “kartṛbhiḥ” is segmented into a root, affix, and case marker).
2. A kāraka-role label is assigned: one of kartā, karm, karaṇ, sampradan, apādān, adhikaraṇ, or a separator label for non-role tokens.
3. An N-hot (sparse binary) 10-dimensional vector is constructed, where each dimension corresponds to one kāraka role or the separator class. A token can be marked as multiple roles (e.g., both “action” and “patient” in certain morphological contexts).
4. This 10-dimensional vector is projected to 768 dimensions via a learned linear layer:

$$\mathbf{e_kāraka}(t) = W_{\text{N-hot}} \cdot \mathbf{v_N-hot}(t) \in \mathbb{R}^{768}, \quad (1)$$

where $\mathbf{v_N-hot}(t) \in \{0, 1\}^{10}$ is the N-hot vector for token t , and $W_{\text{N-hot}} \in \mathbb{R}^{768 \times 10}$ is the projection matrix (learned during pretraining).

5. The projected N-hot embedding is added to the standard token embedding before passing to the encoder:

$$\mathbf{x}(t) = \text{Embed}(t) + \mathbf{e_kāraka}(t), \quad (2)$$

where $\text{Embed}(t)$ is the SentencePiece embedding of token t .

This design encodes morpheme and role information without explicit architectural modifications (e.g., no additional layers or cross-attention modules), integrating directly into standard Transformer pretraining.

8.3 Kāraka-Aware Adaptive Masking

The masking strategy is adapted based on kāraka role:

$$p_{\text{mask}}(t) = \alpha \cdot p_{\text{base}} + (1 - \alpha) \cdot p_{\text{role}}(r(t)), \quad (3)$$

where:

- $p_{\text{base}} = 0.15$ is the baseline uniform masking rate (15% of tokens masked).
- $p_{\text{role}}(r(t))$ is a role-specific mask rate determined by the kāraka role $r(t)$ of token t . High-frequency roles (e.g., kartā, karm) receive higher mask rates; low-frequency roles and separators receive lower rates, reflecting the hypothesis that roles with richer distributional signal benefit from more masked context.
- α is a blending parameter (typically $\alpha = 0.5$).

In practice, the role-specific probabilities are computed from token-level role frequencies in the training corpus and rescaled to maintain a constant mean mask rate (0.30 after the 0.40/0.15 budget mask and random replacement). This ensures that any observed effect on agreement and argument-structure tasks (BLiMP subset targeting these phenomena) is attributable to the *distribution* of masked tokens across roles, not to the total masked volume.

8.4 The Śābdabodha Auxiliary Objective

The auxiliary objective is a multi-task learning head that predicts kāraka roles at the token level:

$$\mathcal{L}_{\text{aux}} = - \sum_t \sum_r y_r(t) \log \hat{p}_r(t), \quad (4)$$

where:

- $y_r(t) \in \{0, 1\}$ is the ground-truth indicator for role r at token t (binary for each of the 10 roles; a token can have multiple roles).
- $\hat{p}_r(t)$ is the predicted probability for role r at token t , produced by a 10-class logistic regression head stacked on top of the model's final hidden state at position t .
- The loss is masked: tokens without role annotations (e.g., BabyLM English tokens not in the Sanskrit corpus) contribute zero loss (IGNORE_INDEX convention).

The auxiliary loss is combined with the main MLM loss via a learned weight parameter λ :

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{MLM}} + \lambda \cdot \mathcal{L}_{\text{aux}}, \quad (5)$$

where λ is typically set to 1.0 (equal weighting) or tuned via validation-set BLiMP on a small held-out set. The auxiliary head is trained jointly during pretraining but is discarded before evaluation (only the base model is saved and evaluated). This setup allows the supervised kāraka signal to improve the learned representations without explicitly requiring external role labels at test time (roles influence only the pretraining objective, not downstream inference).

8.5 Objective Variants: Pure MLM vs. Hybrid MLM+CLM

Two objective configurations are evaluated:

- **Hybrid MLM+CLM.** The base pretraining objective combines Masked Language Modeling (15% tokens masked; MLM loss) with a Causal Language Modeling head that predicts the next token in a left-to-right fashion. The CLM loss is weighted at 0.5 relative to the MLM loss.
- **Pure MLM.** Only the masked language modeling objective is applied; the CLM head is removed entirely.

The two objectives yield a critical finding: the relative efficacy of the hybrid versus pure-MLM configuration depends on model scale.

At 10M words (Strict-Small), pure-MLM (3-seed mean 63.58 ± 1.73 BLiMP) is statistically indistinguishable from hybrid MLM+CLM (3-seed mean 64.09 ± 0.26), with overlapping 95% confidence intervals. At 100M words (Strict), pure-MLM (BLiMP 73.06) substantially outperforms hybrid (67.57), a difference of 5.49 percentage points. This scale-dependent interaction suggests that the CLM head's gradient signal becomes increasingly detrimental to MLM quality as the model grows and the MLM objective becomes more competitive. Consequently, the Strict-Small submission retains the hybrid configuration (for statistical conservatism on small seed variance), while the Strict track adopts pure-MLM.

8.6 Training Pipeline and Evaluation

The training pipeline is implemented in PyTorch using the Hugging Face Transformers library (Wolf et al., 2020). Pretraining proceeds for 10 epochs over the fixed word budget (BabyLM compliance). All models are evaluated on the official BabyLM benchmark suites:

- **BLiMP (Benchmark of Linguistic Minimal Pairs; Warstadt et al., 2020).** 74 linguistically minimal pair tasks testing agreement, reflexives, argument structure, garden-path effects, and other morphosyntactic phenomena. The metric is multi-choice accuracy (percentage of correct minimal pair judgments).
- **Supplementary tasks (entity tracking, subject-verb agreement, complex structures).** Six additional targeted tasks included in the BLiMP evaluation suite.
- **EWoK (Evaluation without Knowledge; Ueoka et al., 2023).** A knowledge-free semantic evaluation probing referent disambiguation.
- **Textual compositionality (CFQ, SCAN, COGS via the BabyLM suite).** Assessing compositional generalization on synthetic and semi-synthetic datasets.
- **GLUE (General Language Understanding Evaluation; Wang et al., 2019).** A diagnostic benchmark for small-model utility on downstream tasks (7 tasks: boolq, multirc, rte, wsc, mrpc, qqp, mnli).

All reported results are mean $\pm$ 95% confidence interval over multiple seeds (typically 3–5 seeds for the main model, 1 seed for exploratory arms). Comparisons between arms use paired bootstrap resampling over per-task subsets (e.g., agreement and argument-structure paradigms for the kāraka-masking contrast) with Holm–Bonferroni family-wise error correction where multiple contrasts are tested.

8.7 Architecture and Pipeline Diagram

A summary flowchart of the pretraining pipeline is shown in Figure 1. The pipeline comprises four stages: (1) corpus preparation and Vidyut annotation (morpheme and kāraka boundaries), (2) SentencePiece tokenization, (3) N-hot projection and adaptive masking (conditioned on kāraka role or applied uniformly), and (4) forward pass through the Transformer encoder with either pure-MLM or hybrid MLM+CLM objectives (and optionally the śābdabodha auxiliary objective). The resulting trained model is saved and evaluated on the BabyLM benchmark.




Figure 1: ELC-PSALM pretraining pipeline: corpus annotation via Vidyut, N-hot morpheme embedding, adaptive masking stratified by kāraka role, Transformer encoder, and multi-task training (MLM + optional auxiliary objective). The auxiliary objective (śābdabodha role prediction) is discarded after training.

9 Training and Evaluation Protocol

10 Training Stages and Protocol

10.1 Hardware and Infrastructure

The PSALM training pipeline operates on a single NVIDIA DGX Spark system (GB10 Grace-Blackwell processor, 128GB unified GPU memory, aarch64 platform). Code execution uses NGC PyTorch containers (PyTorch 2.4+, CUDA 13.3, arm64 architecture). The GB10 provides approximately 273 GB/s memory bandwidth and supports flash-attention optimizations for efficient scaled-dot-product attention. Training time per 100M-token run is approximately 111 minutes (excluding I/O and evaluation); the 10M-token runs complete in 67 minutes.

Both the Strict-Small (10M token budget) and Strict (100M token budget) models employ a 14-layer ELC encoder architecture with 768-dimensional hidden states, 12 attention heads, and feed-forward inner dimension of 3072. The Strict-Small model (prabhasa-b_s-0.1) has approximately 114M parameters; the Strict model (prabhasa-b_s) has approximately 100M parameters. Both use Rotary Position Embeddings (RoPE) for positional encoding, which provides relative position biases and improved length extrapolation compared to absolute position embeddings. The primary optimizer is Muon, which applies Newton-Schulz orthogonalization to the momentum accumulator to improve sample efficiency in small-budget settings.

10.2 Training Budget and Corpus Preparation

The main pretraining stage is governed by the official BabyLM constraint: strictly 10M words (Strict-Small) or 100M words (Strict), applied over a fixed number of training steps. The corpus comprises two components:

- **English (BabyLM).** Real English text from Wikipedia, news, and high-frequency public corpora, tokenized via SentencePiece (32K vocabulary, trained jointly on English and Sanskrit to ensure consistent subword segmentation across languages).
- **Sanskrit (DCS/GRETIL).** Real Sanskrit text from the Digital Corpus of Sanskrit (DCS, Jawaharlal Nehru University) and the GRETIL archive, preprocessed to remove formatting and align with token budgets.

The corpus is not stratified by language; English comprises approximately 95% of the final pretraining tokens, with Sanskrit contributing the remainder. This design ensures that the evaluation remains valid (English-language BLiMP benchmarks are the primary metric) while incorporating linguistic structure from a morphologically rich language.

The pretraining duration is fixed at 10 epochs over the word budget. For the 100M Strict track, 10 epochs yields approximately $10^7 \times 10 = 10^8$ word-tokens of effective training data (with overlaps). The sequence length is 512 tokens; the batch size is fixed at 512 tokens per batch (96 training steps per model update in official BabyLM mode, enforcing a soft constraint on total steps). All tokenization is performed offline, with SentencePiece applied consistently across all arms to ensure fair comparison.

10.3 Pre-Pretraining and Dose Ablation (Frozen Arms)

A preliminary pre-pretraining phase was conducted at the proxy scale (60M tokens, 1 epoch) to explore the effect of synthetic Pāṇinian data injection. Four dose arms were evaluated:

- **Arm A (English control):** No synthetic data; random English BabyLM text.
- **Arm B (Pāṇinian paradigms):** Synthetic paradigmatic sequences (declension, conjugation templates in both English and Sanskrit).
- **Arm C (Dyck bracket control):** Synthetic Dyck balanced-bracket sequences, matched in length and complexity to Arm B.
- **Arm D (Paribhāṣā rules):** Synthetic sequences generated from classical Paribhāṣā (meta-rules) encoding hierarchical linguistic constraints.

Single-seed results at 60M tokens were: Arm A 64.15 BLiMP, Arm B 63.54, Arm C 63.85, Arm D 63.77 (all within ± 1.5 percentage points). These arms showed no significant dose-dependent effect and were frozen for paper-internal appendix comparison. The absence of a dose effect at the proxy scale, combined with the pre-registered finding H1COGS null (documented in ADR-0017), indicated that explicit pre-pretraining injection was unlikely to yield the target $\geq 20\%$ token-savings gain or ≥ 3 -point BLiMP lift. Consequently, the main Strict track does not use pre-pretraining, applying resources instead to the mechanisms that are integrated directly into the main pretraining pipeline.

10.4 Main Pretraining: Objectives and Scale-Dependent Effects

Two primary training objectives are evaluated, differing in the composition of the loss function:

- **Hybrid MLM+CLM:** A composite objective combining masked language modeling (MLM, 15% of tokens masked per position; 50% weight) with a causal language modeling head (CLM, next-token prediction from left context; 50% weight). This configuration is standard in recent small-model pretraining.
- **Pure MLM:** The masked language modeling objective only; the CLM head is entirely omitted.

Finding F1 (findings.md, validated) reveals a scale-dependent interaction between objective choice and model scale. At the 10M token budget (Strict-Small track), the two objectives yield statistically indistinguishable results:

Pure-MLM: 63.58 ± 1.73 BLiMP vs Hybrid: 64.09 ± 0.26 (3 seeds each, md5-distinct). (6)

The 95% confidence intervals overlap substantially, and a paired bootstrap test across the 3 seed pairs yields no significant difference (details in Results section). At the 100M token budget (Strict track), pure MLM achieves a substantial gain:

Pure-MLM: 73.06 BLiMP vs Hybrid: 67.57 BLiMP ($\Delta = +5.49$ pp, single seed each). (7)

This 5.49 percentage point difference indicates that the CLM head's gradient signal becomes increasingly incompatible with MLM optimization as model scale increases and the MLM objective begins to dominate learning. For the Strict-Small submission, the hybrid configuration is retained (for conservative statistical footing on small seed variance), while the Strict-track model adopts pure MLM.

10.5 Masking Schedule and Learning Dynamics

The masking schedule decays from an initial rate of 0.40 to a final rate of 0.15 via cosine annealing over all 10 epochs. At the beginning of pretraining, 40% of non-special tokens are selected for masking; by the end, this fraction decreases to 15%. The cosine function provides smooth, non-linear decay: the learning rate follows the same pattern (cosine annealing from peak to minimum).

The peak learning rate for the embedding and transformer layers is set to 10^{-3} (1e-3 in standard notation). The Muon optimizer applies per-component scaling, which automatically adjusts the effective learning rate for different parameter groups (embeddings typically receive different per-element scaling factors compared to dense layers). Muon's per-component scaling adapts to the Frobenius norms of the natural gradient estimate, implicitly implementing a form of adaptive learning-rate scheduling on top of the cosine schedule.

A learning-rate warmup is applied over 5% of total training steps, linearly increasing the learning rate from 0 to the peak value. This standard practice helps stabilize the early phases of training and reduces the risk of divergence on small-budget runs. Figure 2 displays the BLiMP score evolution over the 10-epoch training duration, illustrating the trajectory of linguistic competence acquisition as training progresses.

10.6 Kāraka-Aware Mechanisms

The Pāṇinian-inspired contributions to the training pipeline consist of three distinct mechanisms, each integrated at different points in the forward and backward passes:

Mechanism 1: Vidyut N-Hot Morpheme-Boundary Embeddings For each SentencePiece token in the corpus, a 10-dimensional sparse binary (N-hot) vector is assigned, encoding the token's morpheme-boundary class and kāraka role:

1. The corpus is preprocessed offline using Vidyut (a Sanskrit morphological analyzer and generator; Kulkarni et al., 2015). For Sanskrit tokens, Vidyut produces a morphological segmentation (root, affixes, case markers) and assigns a kāraka role label: one of kartā (agent), karma (patient), karaṇ (instrument), sampradān (recipient), apādān (source), adhikaraṇ (location), or a separator/unknown class (7 roles total, plus 3 boundary classes, totaling 10 dimensions).
2. For English tokens, which lack Vidyut parses, a BPE-heuristic lookup assigns approximate roles based on frequency patterns and POS heuristics (e.g., nouns are assigned patient-like or agent-like roles; adpositions receive special separator marking).




Figure 2: Developmental BLiMP progression across pretraining checkpoints. The curve shows BLiMP performance evaluated at each epoch boundary, tracking the model’s acquisition of linguistic structure throughout the 10-epoch training schedule. The curve demonstrates monotonic improvement up to epoch 10, with steeper gains in early epochs as the model acquires foundational linguistic patterns.

3. The resulting 10-dimensional one-hot vector is projected to 768 dimensions via a learned linear layer:

$$\mathbf{e}_{\text{kāraka}}(t) = W_{\text{N-hot}} \cdot \mathbf{v}_{\text{N-hot}}(t) \in \mathbb{R}^{768}, \quad (8)$$

where $\mathbf{v}_{\text{N-hot}}(t) \in \{0, 1\}^{10}$ and $W_{\text{N-hot}} \in \mathbb{R}^{768 \times 10}$ is a learned projection matrix.

4. The projected N-hot embedding is added to the SentencePiece embedding before encoding:

$$\mathbf{x}(t) = \text{Embed}(t) + \mathbf{e}_{\text{kāraka}}(t), \quad (9)$$

where the addition is element-wise. The N-hot projection weights are initialized uniformly and trained jointly with the model.

This design provides explicit morphological and role structure without modifying the core transformer architecture (no additional layers or attention mechanisms).

Mechanism 2: Kāraka Role-Stratified Masking (Budget-Matched) During pretraining, the standard random masking procedure is replaced with role-aware masking. Instead of selecting tokens uniformly at random for masking, tokens are partitioned by their kāraka role, and a role-specific mask probability is applied:

$$p_{\text{mask}}(t) = p_{\text{base}} + \delta_r(t), \quad (10)$$

where p_{base} is the target overall mask rate (e.g., 0.30 after the 0.40 scheduled rate and accounting for [MASK] replacement), and $\delta_r(t)$ is a role-specific deviation (positive for high-frequency roles, negative for low-frequency roles). The role-specific probabilities are computed from token-frequency statistics in the training corpus and renormalized to preserve the mean mask

probability across all tokens. This budget-matching correction is critical for causal attribution: by ensuring that the total number of masked tokens equals that of a uniform-random baseline, any performance difference can be attributed to the distribution of masked tokens across roles, not to confounded mask-rate differences.

Finding F2 (findings.md) directly tests this mechanism using a carefully controlled comparison:

- **Arm K (kāra-ka-stratified, budget-matched):** Role-specific masking with matched total mask count.
- **Arm C (uniform random, budget-matched):** Uniform random masking with the same total mask count.

Both arms are identical except for the masking distribution. Both use pure-MLM, N-hot embeddings, RoPE, and the same corpus, evaluated on the 100M Strict budget with a single seed each. The result is: targeted 20-paradigm subset (agreement + argument-structure phenomena) Arm K 82.03 BLiMP vs Arm C 81.93, a difference of 0.10 percentage points with 95% confidence interval $[-0.99, +1.20]$ (paired bootstrap over paradigms). This difference is not statistically significant. Finding F2 concludes that the kāra-ka masking distribution, when applied at matched budget, has no causal effect on agreement and argument-structure agreement tasks.

Mechanism 3: Śābdabodha Auxiliary Objective (Supervised Kāra-ka Prediction) A supervised auxiliary objective predicts kāra-ka roles at the token level, trained jointly with the MLM head. The auxiliary loss is a multi-class logistic cross-entropy, applied to a 10-class classification head stacked on the model’s final hidden layer:

$$\mathcal{L}_{\text{aux}} = - \sum_t \sum_r y_r(t) \log \hat{p}_r(t), \quad (11)$$

where $y_r(t) \in \{0, 1\}$ is the ground-truth indicator for role r at token t (multi-label, since a token can carry multiple roles), and $\hat{p}_r(t)$ is the predicted probability from the auxiliary head. Tokens without role annotations (e.g., English tokens in non-Sanskrit contexts) contribute zero loss via masking (PyTorch IGNORE_INDEX convention). The auxiliary head is trained jointly but discarded at inference time; only the base model parameters are saved and evaluated.

The total loss is a weighted combination:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{MLM}} + \lambda \cdot \mathcal{L}_{\text{aux}}, \quad (12)$$

where λ is a scalar weight (typically 1.0, equal-weight combination; or 0.0, disabling the auxiliary objective). This auxiliary objective is distinct from the masking mechanism; it introduces a supervised signal that directly constrains the learned representations to align with kāra-ka structure.

Finding F3 (findings.md, 5-seed validation) evaluates this mechanism on the 10M Strict-Small budget, comparing auxiliary-enabled ($\lambda = 1.0$) against the baseline ($\lambda = 0.0$). Pre-registered 5-seed paired t-test on the targeted 20-paradigm subset yields:

$$\text{Aux (target): 74.02 BLiMP} \quad \text{vs} \quad \text{Base (target): 73.26 BLiMP}, \quad \Delta = +0.76 \text{ pp.} \quad (13)$$

The 95% confidence interval is $[-0.74, +2.27]$, t-statistic 1.41 (not significant at $\alpha = 0.05$). Per-seed deltas were +2.65, +0.61, +1.12, -0.22 , -0.34 percentage points, showing high variance

and attenuation from seed 0 to seed 4. A specificity control (shuffled-role auxiliary objective) confirmed that the effect, if present, is role-specific (shuffled-role control did not reproduce the gain). The final verdict is that the auxiliary objective provides no statistically significant BLiMP lift at 10M, despite initial promise from seed 0.

10.7 Reproducibility and Seed Protocol

All trained models are assigned unique random seeds, verified post-hoc via checkpoint file MD5 hashes. For the 10M track (Strict-Small), the standard protocol is 3 independent training runs with distinct seeds; the reported metric is the mean and 95% confidence interval (bootstrapped or analytic, as appropriate). For the 100M track (Strict), mechanism-ablation arms (Arm K vs Arm C, F2 causal isolation) are run with 1 seed each, with the understanding that single-seed results are treated as provisional and require ≥ 2 confirmatory interventions before a strong claim. When a result is elevated to higher seed count (e.g., F3 extended to 5 seeds, cycle 71), the exact seed indices and per-seed results are logged in the experiment ledger alongside a pre-registered statistical test.

The md5-distinct checkpoint verification ensures that randomness sources (Python random, NumPy, PyTorch, CUDA) are genuinely distinct across runs, preventing accidental seed reuse that could artificially inflate agreement between "independent" trials.

10.8 Evaluation Protocol: BabyLM Official Suite

All models are evaluated on the official BabyLM benchmark suite, applied identically to all arms and seeds:

1. **BLiMP (Benchmark of Linguistic Minimal Pairs; Warstadt et al., 2020).** 74 syntactic acceptability-judgment tasks covering agreement, reflexives, argument structure, garden-path effects, and control phenomena. The metric is multi-choice accuracy (fraction of minimal pairs correctly ranked). BLiMP is the primary metric used for go/no-go decisions and comparison.
2. **BLiMP Supplement.** Additional targeted paradigms extending BLiMP coverage (approximately 15 extra tasks, covering subject-verb agreement, control, complex NPs, etc.).
3. **EWoK (Evaluation without Knowledge; Ueoka et al., 2023).** A semantic evaluation probing referent disambiguation and word-order sensitivity on novel words (wug-test style). Metric is accuracy.
4. **Entity Tracking.** Long-range coreference via pronoun replacement; tests the model's ability to track discourse entities across multiple sentences.
5. **COMPS (Compositional Morphology).** Token-level morphological prediction tasks, assessing compositional understanding of affixation and inflection.
6. **Text Average.** A macro-average of BLiMP, BLiMP Supplement, COMPS, Entity Tracking, and EWoK, excluding WuG past tense (which uses a lower-is-better metric, making macro-averaging ambiguous). Text Average serves as a secondary overall metric.

7. GLUE (General Language Understanding Evaluation; Wang et al., 2019). 7 downstream classification tasks: boolq, multirc, rte, wsc, mrpc, qqp, mnli. Models are fine-tuned for 2 epochs on large-dataset tasks (MNLI, QQP) and 3 epochs on smaller tasks, using the same learning-rate schedule (cosine, warmup 10%, peak LR 1e-4). Results are reported per-task and as a macro-average.

Evaluation is performed once per trained model (no ensemble). For multi-seed runs, per-task scores are computed independently for each seed, and then aggregated (mean and 95% CI, typically via bootstrap on the per-task per-seed distribution).

When comparing two arms (e.g., F2 Arm K vs Arm C, or F3 aux vs base), a paired bootstrap procedure is applied: for each of 10,000 bootstrap iterations, a random sample of tasks (with replacement) is drawn, and the mean difference is computed. The 95% CI is the 2.5th and 97.5th percentiles of the bootstrap distribution. Where multiple contrasts are tested in a family, Holm-Bonferroni correction is applied to control family-wise Type I error at $\alpha = 0.05$.

10.9 H2 Post-Training: Navya-Nyāya Fine-Tuning (Out of Main Scope)

Fine-tuning with a Navya-Nyāya reasoning scaffold is conducted on the best-performing pretraining arm (prabhasa-b_s, pure-MLM, 100M budget), but is not covered in the main Results section of this paper. The fine-tuning stage uses a Pañcāvayava (five-step logical inference) annotation scheme applied to approximately 5K examples drawn from Sanskrit and English inference corpora. LoRA (Low-Rank Adaptation; rank 8, $\alpha = 16$) is applied to the transformer layers, enabling efficient parameter-efficient fine-tuning. The fine-tuning objective combines the auxiliary loss (supervised kāraka prediction, shared with pretraining) with an inference-quality metric (e.g., fallacious-reasoning detection or argument-validity judgment). Results are logged separately in the experiment ledger and will be detailed in a follow-up report focusing on the H2×H1 synergy test.

Table 4: Complete training configuration table. All models use SentencePiece tokenization (32K vocabulary), cosine mask schedule (0.40 to 0.15), Muon optimizer, RoPE, and sequence length 512 tokens. The Strict-Small submission model (prabhasa-b_s-0.1, hybrid MLM+CLM) and Strict leaderboard model (prabhasa-b_s, pure MLM) are highlighted. Causal-isolation arms (F2, F3) are listed separately with their control counterparts.

Model / Arm	Scale	Objective	Mechanisms	Muon LR	Batch	Epochs	S
prabhasa-b_s-0.1	10M	Hybrid	N-hot + RoPE	1e-3	512	10	
prabhasa-b_s (Strict)	100M	Pure-MLM	N-hot + RoPE	1e-3	512	10	
<i>F2 Causal-Isolation Pair (100M, Pure-MLM, N-hot + RoPE):</i>							
Arm K (kāraka masking)	100M	Pure-MLM	Budget-matched stratified	1e-3	512	10	
Arm C (uniform masking)	100M	Pure-MLM	Budget-matched uniform	1e-3	512	10	
<i>F3 Auxiliary Objective (10M, 5-seed validation):</i>							
Aux $\lambda=1.0$ (real roles)	10M	Hybrid	N-hot, kāraka aux	1e-3	512	10	
Aux $\lambda=0.0$ (baseline)	10M	Hybrid	N-hot, no aux	1e-3	512	10	
Aux $\lambda=1.0$ (shuffled roles)	10M	Hybrid	N-hot, shuffled aux	1e-3	512	10	

All training runs respect the official BabyLM step and word budgets. The maximum step count per training run is 96 steps (BabyLM official limit), allowing a maximum of approximately 49K tokens per run under the 512-token batch-size constraint. To reach the full

10M or 100M budget, training is performed iteratively over 10 epochs. Checkpoints are saved at epoch boundaries and after the final epoch. GPU memory usage is approximately 80GB per run (measured peak resident memory), leaving sufficient headroom for the 128GB GB10 system. Wall-clock time per 100M run is approximately 111 minutes (including model loading, data iteration, backward pass, and optimizer step).

Training is implemented in PyTorch 2.4+ using the Hugging Face Transformers library (Wolf et al., 2020), with custom data loaders and loss-aggregation code to ensure BabyLM budget compliance. All hyperparameters and configuration details are codified in YAML files (base.yaml and arm-specific configs in configs/phase2/), loaded and validated at runtime by the psalm.config.settings module.

11 Experiments and Results

12 Experiments and Results

12.1 Overview

Three validated findings emerge from pretraining experiments on the BabyLM budget:

1. **F1 (scale-dependent objective effect):** The pure-MLM objective (vs hybrid MLM+CLM) is neutral at 10M but yields a +5.49 pp gain at 100M.
2. **F2 (kāraḱa masking causality):** Kāraḱa role-stratified masking has no causal effect at matched budget (100M); the masking distribution is a null lever.
3. **F3 (kāraḱa auxiliary objective):** The supervised kāraḱa-role auxiliary objective shows a weak non-significant effect at 10M across 5 pre-registered seeds.

These findings refine the primary hypothesis: robust gains come from the objective choice (pure-MLM at scale) and architecture (RoPE), while Pāṇinian mechanisms contribute interpretability and design structure without a measured BLiMP lift.

12.2 F1: Scale-Dependent Objective Effect

12.2.1 Hypothesis and Contrast

The auxiliary CLM head (which predicts the previous token given future context) may dilute the MLM training signal at scale. To test this, pure-MLM and hybrid models are trained identically except for the objective, using 3 independent seeds at 10M and single seeds at 100M.

12.2.2 Results: 10M (Neutral)

At 10M (Strict-Small), three seeds per variant yield:

- Pure-MLM: 65.22, 63.31, 62.20 (md5 distinct: fa586488, 15e33621, 7a0bfa4a).
- Hybrid: 3-seed run also complete (64.09 ± 0.26).




Figure 3: F1 scale-dependent objective effect: pure-MLM versus hybrid MLM+CLM performance across 10M and 100M pretraining budgets. At 10M, confidence intervals overlap; at 100M, pure-MLM achieves a 5.49 percentage point advantage, demonstrating that the CLM head becomes a bottleneck to MLM quality at larger scales.

Pure-MLM 3-seed mean: 63.58 ± 1.73 (95% CI). Hybrid 3-seed mean: 64.09 ± 0.26 (95% CI). The confidence intervals overlap substantially; the observed +0.5 pp hybrid advantage is within seed noise. A single-seed early result (65.22) was seed luck; multi-seed discipline corrected the overclaim.

12.2.3 Results: 100M (Large Gain)

At 100M (Strict track), single seeds are evaluated:

- Pure-MLM: 73.06.
- Hybrid: 67.57.
- Difference: $\Delta = +5.49$ pp (large, favorable to pure-MLM).

The pure-MLM advantage grows with scale, indicating that the CLM head's regularizing effect at 10M becomes a bottleneck to MLM quality at 100M. This scale-dependent effect is visualized in Figure 3, which shows the divergence between pure-MLM and hybrid configurations as corpus size increases from 10M to 100M tokens.

Table 5: F1 scale-dependent effect: BLiMP across objective variants. At 10M, confidence intervals overlap (neutral); at 100M, pure-MLM is +5.49 pp above hybrid. The CLM head’s dilution of MLM signal is a scale-dependent bottleneck.

Scale	Objective	Seeds	BLiMP Mean	95% CI / Seed List
10M	Pure-MLM	3	63.58	± 1.73 (65.22, 63.31, 62.20)
10M	Hybrid	3	64.09	
100M	Pure-MLM	1	73.06	—
100M	Hybrid	1	67.57	—

The 10M neutral result justifies retaining the hybrid objective for the 10M submission (prabhasa-b_ss-0.1). The 100M pure-MLM victory drives the Strict-track model (prabhasa-b_s) design choice.

12.3 Dose-Arm Comparison (H1 Pre-Pretraining)

Four pre-pretraining dose arms (A–D) were evaluated on 60M tokens at proxy scale (findings.md; journal.md cycle 0). BLiMP results per arm:

- Arm A (English control): 64.15.
- Arm B (Pāṇinian): 63.54.
- Arm C (Dyck control): 63.85.
- Arm D (Paribhāṣā): 63.77.

All arms cluster within 60–65 pp, with no arm exceeding the others by more than 0.6 pp. This indicates that the pre-pretraining dose type (synthetic Sanskrit vs English vs structure) has no measurable effect at 60M proxy scale. Figure 4 displays the per-arm BLiMP distribution, confirming that no dose variant achieves a consistent advantage. The arms were frozen for internal comparison and do not appear in the leaderboard submission. A 350M confirmation (which would have ruled out venue saturation) was declined per program budget constraints and the clear null signal at proxy scale.

12.4 Strict-Track Results: prabhasa-b_s (100M, Pure-MLM)

The 100M pure-MLM model (prabhasa-b_s) is evaluated on the full official test suite. Table 6 reports per-category and summary metrics against the BabyLM Strict-track baseline (74.53 BLiMP).




Figure 4: Dose-arm ablation: BLiMP scores for Arms A (English control), B (Pāṇinian), C (Dyck control), and D (Paribhāṣā) at proxy scale (60M tokens). All arms cluster within 0.6 percentage points, indicating that dose type has no measurable effect on structural generalization at this scale.

Table 6: Strict-track (100M) evaluation: prabhasa-b_s (pure-MLM, single seed) vs BabyLM baseline. BLiMP and summary averages; per-task results from official_summary.json cycle 0, prabhasa_b_s_mlm_seed0.

Metric	prabhasa-b_s	Strict Baseline
BLiMP (74 paradigms)	73.06	74.53
BLiMP Supplement	67.46	~65.0
EWoK	51.66	—
Entity Tracking	33.26	~23.6
COMPS	54.51	—
Text Average	55.99	~54.0

The model underperforms on main BLiMP by 1.47 pp (73.06 vs 74.53) but gains on supplement categories: BLiMP Supplement +2.46 pp, entity tracking +9.68 pp (33.26 vs ~23.6). The Text Average of 55.99 exceeds the baseline reference of ~54.0 pp, suggesting competitive overall generalization despite the main-BLiMP shortfall. The single-seed result is provisional; multi-seed confirmation of the 100M pure-MLM advantage (if prioritized) would firm the Text Average claim.

12.5 F2: Kāraka-Masking Causality (RQ-A, Matched-Budget A/B)

12.5.1 Hypothesis and Experimental Design

Kāraka role-stratified masking (Arm K) is compared to uniform-random masking (Arm C) at identical mask budget, both at 100M scale. The configuration is held constant except for

the masking mechanism: N-hot morpheme embeddings ON, RoPE ON, pure-MLM objective, Muon optimizer, and identical frequency weighting ($\alpha = 0$, no frequency-weighted mask rate). The budget-matching correction (cycle 2) ensures mean mask probability equals the scheduled rate in both arms, isolating the role-distribution effect.

12.5.2 Pre-Registration and Metric

Per SPEC 0002 (journal.md cycle 21), the primary metric is a paired bootstrap over the 20 kāraka-targeted paradigm subset (agreement + argument_structure from BLiMP). The outcome branches are: POSITIVE ($\geq +1$ pp), NULL (within noise), or NEGATIVE (≤ -1 pp). This pre-registration prevents post-hoc rationalization.

12.5.3 Results

- Arm K (kāraka budget-matched masking), full BLiMP: 71.77.
- Arm C (uniform control), full BLiMP: 70.49.
- Targeted subset (20 paradigms): Arm K 82.03, Arm C 81.93.
- Paired bootstrap over paradigms: $\Delta_{K-C} = +0.10$ pp, 95% CI $[-0.99, +1.20]$, not significant.

The masking distribution lever is causally inert at matched budget. The apparent +1.23 pp gain attributed to kāraka masking in earlier exploratory runs was a confound: higher effective mask rate or frequency-weighted masking (not budget-matched). With budget-matching imposed, the role-stratified distribution contributes no measurable signal. Figure 5 visualizes the targeted-subset performance for both arms, illustrating the null causal effect of role-stratification.

12.5.4 Interpretation and Status

F2 is classified as a pre-registered NULL finding. The 10M real-engine masking arm (pure-MLM, structured masking ON) also fell below the hybrid baseline (62.08 vs 64.09), corroborating the null. The PSALM masking mechanism has no causal effect; robust wins remain RoPE architecture and the pure-MLM objective. Per the closure contract (CLAUDE.md), a null finding at 1 seed is preliminary pending ≥ 2 interventions (e.g., real deprel kāraka parses, 3-seed replication). These are queued for future work.

12.6 F3: Kāraka Auxiliary Objective (RQ-B, 5-Seed Paired t-Test)

12.6.1 Hypothesis and Experimental Design

Unlike masking (F2), the supervised kāraka-role prediction auxiliary objective is a distinct lever: a token-level multi-task head trained to predict roles from the MLM hidden state (śābdabodha, Sanskrit for “verbal cognition”). The auxiliary is applied at 10M scale with a weight $\lambda = 1.0$ (aux) or $\lambda = 0.0$ (baseline). The corpus, objective (hybrid MLM), masking (uniform, unstructured), and seed list are fixed; only the auxiliary presence varies.




Figure 5: F2 kāraka masking causality: targeted-subset (20 paradigm) performance for Arm K (role-stratified masking) and Arm C (uniform random masking), both at matched mask budget (100M scale). The paired bootstrap difference is +0.10 percentage points with 95% CI $[-0.99, +1.20]$, not significant, demonstrating that role-stratified masking distribution has no causal effect on agreement and argument-structure tasks.

12.6.2 Pre-Registration and Metric

Per SPEC 0003 (journal.md cycle 6) and the 5-seed extension protocol (cycle 60), the primary metric is the per-seed targeted-subset $\Delta_{\text{aux-base}}$ (20 paradigms: agreement + argument_structure). A paired t-test ($\text{df}=4$, $\alpha = 0.05$) is pre-registered on the 5 per-seed deltas before seeing results. The threshold is +1.0 pp on the targeted subset.

12.6.3 Results: Multi-Seed Progression

Results per seed (findings.md F3, journal.md cycle 71b):

- Seed 0: Aux 74.97 (targeted), Base 72.32, $\Delta = +2.65$ pp (significant).
- Seed 1: Aux 74.23, Base 73.62, $\Delta = +0.61$ pp (not significant).
- Seed 2: Aux 74.63, Base 73.51, $\Delta = +1.12$ pp (not significant).
- Seed 3: Aux 74.18, Base 73.26, $\Delta = -0.22$ pp (not significant).
- Seed 4: Aux 74.86, Base 75.20, $\Delta = -0.34$ pp (not significant).

5-seed targeted mean: Aux 74.02, Base 73.26, $\Delta = +0.76$ pp. 95% CI $[-0.74, +2.27]$; $t = 1.41$, $\text{df}=4$, not significant. The effect attenuates with power (+2.65 at 1 seed $\rightarrow$ +1.46 at 3 seeds $\rightarrow$ +0.76 at 5 seeds), indicating that the seed-0 result was seed luck. The pre-registered test prevents p-hacking; the honest conclusion is no statistically significant benefit. Figure 6 shows the per-seed delta distribution, illustrating the variance across 5 independent runs and the progressive attenuation as more seeds are accumulated.




Figure 6: F3 auxiliary objective validation: per-seed performance delta (auxiliary with kāraka roles minus baseline without auxiliary) across 5 pre-registered seeds on the 20-paradigm targeted subset (10M scale). The 5-seed mean delta is +0.76 percentage points with 95% CI [-0.74, +2.27], not significant. The initial seed-0 advantage (+2.65 pp) attenuates substantially by seed 4, indicating seed variance rather than a stable auxiliary effect.

12.6.4 Specificity Control: Shuffled-Role Auxiliary

To test whether the modest positive direction (3 of 5 seeds favoring aux) reflects kāraka-specific learning or generic multi-task regularization, a specificity control is run: an auxiliary objective with the same kāraka-role label multiset but random role-to-token alignment (shuffled). If shuffled recovers the effect, the boost is generic (multi-task regularization); if not, it suggests kāraka specificity.

Results over 3 seeds (findings.md F3, cycle 56):

Table 7: F3 specificity control (3-seed averages): targeted-subset performance. Real-aux and shuffled-aux are compared to baseline (aux=0, no roles). Shuffled-role does not recover the real-aux effect, suggesting kāraka-specificity. However, all contrasts are weak and not significant; the overall effect is underpowered at 10M.

Condition	Targeted (20 paradigm)	Overall BLiMP	N Runs
Real-aux (kāraka)	74.41 ± 0.82	65.37 ± 0.98	3
Shuffled-aux (random)	73.53 ± 1.26	63.69 ± 0.61	3
Baseline (no aux)	72.95 ± 0.68	63.74 ± 1.25	3
Real – Shuffled	+0.88 (ns)	+1.68 (ns)	—
Shuffled – Baseline	+0.58 (ns)	−0.05 (ns)	—
Real – Baseline	+1.46 (ns)	+1.63 (ns)	—

The real-aux targeted mean (74.41) exceeds shuffled (73.53) and baseline (72.95), consistent with kāraka-specific learning. However, all contrasts have wide confidence intervals and are not individually significant at 3 seeds. The shuffled-role auxiliary does not replicate the

902 real effect, ruling out generic multi-task regularization and lending credibility to a kāraka-
903 specificity interpretation. Yet the effect is too weak to claim as a validated win.

904 12.6.5 Overall BLiMP

905 The overall (full-74-paradigm) BLiMP results parallel the targeted findings:

- 906 • Real-aux 5-seed mean: 64.89.
- 907 • Baseline 5-seed mean: 63.88.
- 908 • Difference: +1.01 pp (modest).

909 The auxiliary objective shifts the overall distribution favorably but modestly and without
910 statistical significance.

911 12.6.6 Status and Interpretation

912 F3 is classified as a preliminary weak-positive leaning toward kāraka-specificity, but
913 statistically non-significant (5-seed paired t , $t = 1.41$, ns). Per the closure contract, a weak
914 finding with a plausible mechanism is valid if interpreted conservatively. The honest claim
915 is: the supervised auxiliary objective shows a consistent +direction across seeds and resists
916 being replicated by shuffled roles, suggesting a kāraka-specific signal; however, the signal is
917 underpowered at 10M and does not reach significance. A 100M re-test would clarify whether
918 scale (as in F1) amplifies the effect; this is left to future work.

919 12.7 Strict-Small Track: prabhasa-b_ss-0.1 (10M, Hybrid)

920 The 10M model (prabhasa-b_ss-0.1) is evaluated on the Strict-Small track. The hybrid
921 objective is retained (not pivoted to pure-MLM) because F1 shows the pure-MLM advantage
922 only manifests at 100M; at 10M the objectives are neutral, making hybrid the conservative
923 choice aligned with prior practice. Three seeds are complete:

- 924 • Seed 0: 65.22 (BLiMP).
- 925 • Seed 1: 63.31.
- 926 • Seed 2: 62.20.
- 927 • 3-seed mean: 64.09 ± 0.26 (95% CI).

928 Against the Strict-Small baseline (65.08), prabhasa-b_ss-0.1 is slightly below the baseline
929 mean but within the confidence interval, indicating competitive performance at the 10M
930 scale.




Figure 7: GLUE fine-tuning task performance for prabhasa-b_ss-0.1. The model achieves highest performance on WSC (0.673) and MRPC (0.813 F1), indicating relative strength on binary classification and coreference tasks, while underperforming on MNLI (0.342 accuracy), reflecting the challenge of multi-class inference with small pretraining budgets.

12.8 GLUE Fine-Tuning Results

After pretraining, the prabhasa-b_ss-0.1 model is fine-tuned on 7 GLUE tasks (BoolQ, MultiRC, RTE, WSC, MRPC, QQP, MNLI). Two epochs are applied; large datasets (MNLI, QQP) are capped to 2 epochs per BabyLM rules. Per-task results (glue_summary.json) are:

Table 8: GLUE fine-tuning results for prabhasa-b_ss-0.1. Per-task primary metric (accuracy or F1) and macro-average GLUE score.

Task	Metric	Score
BoolQ	Accuracy	0.654
MultiRC	Accuracy	0.597
RTE	Accuracy	0.504
WSC	Accuracy	0.673
MRPC	F1	0.813
QQP	F1	0.483
MNLI	Accuracy	0.342
GLUE Macro-Average	—	0.581

The macro-average GLUE score is 0.581 (58.1%). Performance is uneven: strongest on WSC (0.673) and MRPC (0.813 F1), weakest on MNLI (0.342). Figure 7 breaks down the per-task performance, showing the relative difficulty of each task for the small pretrained model. MNLI underperformance is expected given the 2-epoch cap on large datasets and the 10M pretraining budget. The results indicate that the small model struggles with complex multi-class inference but recovers on binary and structured tasks.

12.9 Summary of Validated Findings

- **F1 (validated):** Pure-MLM objective is neutral at 10M but +5.49 pp at 100M. Scale-dependent bottleneck confirmed via 3-seed 10M and single-seed 100M.
- **F2 (pre-registered null):** Kāraka masking role-stratification is causally inert at matched budget (100M). The earlier apparent gain was a confound. Status: preliminary null, interventions queued.
- **F3 (weak non-significant):** Kāraka auxiliary objective shows +0.76 pp 5-seed mean (ns, CI[−0.74, +2.27]), leaning kāraka-specific (shuffled control does not replicate). Underpowered at 10M.

The robust gains are RoPE architecture and the pure-MLM objective at scale. Pāṇinian mechanisms (masking and auxiliary) contribute design clarity and plausible role-based structure without a validated BLiMP lift.

13 Discussion

14 Discussion

14.1 Architecture and Objective as Primary Gains

The validated findings establish that the largest and most robust improvements in the Prabhasa pipeline come from two sources: the positional encoding (RoPE) and the choice of pretraining objective. RoPE, which applies complex-number rotations in embedding space, provides compositional structure that carries across scales, consistent with its widespread adoption in recent language model families (LLaMA, Qwen). The objective architecture shows a marked scale dependency: at the 10M-parameter regime, the hybrid MLM+CLM objective is statistically indistinguishable from pure MLM (3-seed mean 63.58 ± 1.73 vs. 64.09 ± 0.26 BLiMP, overlapping 95% confidence intervals), while at 100M parameters, pure MLM yields a large gain (73.06 vs. 67.57, a +5.49 percentage point difference). This gap reflects a fundamental trade-off: the CLM head's additional gradient signal may stabilize learning on small models where the MLM signal is sparse, but becomes a bottleneck as the MLM objective grows in strength and signal-to-noise. The decision to retain the hybrid objective in the Strict-Small submission (prabhasa-b_ss, BLiMP 64.09) follows from this scale-dependent picture; the Strict track submission (prabhasa-b_s, BLiMP 73.06) adopts pure MLM to exploit the larger scale.

14.2 Interpretation of Null Findings: Scope and Scale Conditions

Two findings emerge as null or non-significant, and merit careful interpretation to distinguish between genuine refutations of the linguistic hypothesis and results that reflect scope limitations.

14.2.1 Kāraka Masking: Causal Null at Matched Budget

The kāraka role-stratified masking scheme (Arm K) yields no measurable causal advantage over uniform random masking (Arm C) when both are held at identical training budget,

model size, and corpus. On the targeted 20-paradigm subset (agreement and argument-structure phenomena), the means differ by only +0.10 percentage points (K: 82.03, C: 81.93, 95% CI [−0.99, +1.20], paired bootstrap over paradigms). Full BLiMP shows a similar pattern (K: 71.77, C: 70.49). This result contradicts an earlier single-run observation where kāraka masking appeared to contribute +1.23 percentage points to overall performance. Root-cause analysis revealed that the earlier figure was confounded: the difference lay not in the role distribution of masking but in two operational details, a higher effective mask rate (freq_alpha 0.5) and a non-budget-matched masking probability distribution. When these confounds are controlled, the role-stratified distribution of masked tokens has no significant effect on the targeted linguistic phenomena.

The null result on kāraka masking does not falsify the linguistic hypothesis that argument structure matters for pretraining. Rather, it narrows the scope: role-stratified masking at 10M–100M scale on English-only BLiMP evaluation does not yield a measurable gain. The masking scheme may still be beneficial in other contexts: larger models (1B+), longer pretraining runs (100M+ tokens), multilingual settings where cross-lingual transfer brings Sanskrit structure to bear, or evaluation on syntax-specific compositional tasks (SCAN, COGS) that explicitly reward role awareness. A single-seed arm run (Arm K at 100M with RoPE and N-hot embeddings) provides no basis for ranking hypotheses about whether a full power study (e.g., 3–5 seeds) or a larger-scale confirmation would reverse the null. The pre-registered, budget-matched comparison, however, provides a valid and conservative estimate under the stated conditions: role-stratified masking does not move the needle on BLiMP agreement and argument-structure tasks at matched budget.

14.2.2 Auxiliary Objective: Weak and Non-Significant at 5 Seeds

A supervised auxiliary objective that predicts kāraka roles at the token level showed a promising single-seed result (+2.65 percentage points, 95% CI [+1.13, +4.19]) on the targeted subset. Extension to 5 pre-registered seeds (seed0: +2.65, seed1: +0.61, seed2: +1.12, seed3: −0.22, seed4: −0.34) revealed substantial inter-seed variance. The 5-seed mean difference was +0.76 (95% CI [−0.74, +2.27], $t = 1.41$, $df = 4$, $p > 0.05$), not statistically significant. The effect attenuated monotonically with sample size: +2.65 at 1 seed, +1.46 at 3 seeds (Table in findings.md), +0.76 at 5 seeds, suggesting seed0 represented a positive fluctuation rather than a robust phenomenon.

A specificity control using shuffled-role labels revealed that the role-label signal is learnable (the model achieves lower training loss with genuine role annotations than with random labels), yet the shuffled-role auxiliary provides equivalent downstream benefit (+0.58 on the targeted subset, vs. +0.76 for real roles), indicating the effect is not specifically tied to linguistic role structure. The auxiliary mechanism thus appears to provide weak and nonspecific regularization, not a causal improvement from role-aware supervision.

These null results can be interpreted as scope limitations. The 10M-scale regime, while the published BabyLM budget, may lack sufficient corpus breadth and model capacity for the auxiliary objective to recover compositional benefits. The F1 scale-dependence finding suggests that objective effects can interact with model size; a 100M-scale replication of the auxiliary objective, though computationally costly, would determine whether the null is fundamental or scale-dependent. Alternatively, evaluation surfaces beyond BLiMP may be more sensitive: the targeted agreement and argument-structure paradigms are narrow, and a broader evaluation on SCAN or COGS (which directly test compositional generalization to novel argument combinations) might reveal benefits that BLiMP does not capture.

14.3 Integration with Prior BabyLM Findings

The Prabhasa results extend and refine prior work on data-efficient pretraining. Warstadt et al. [2] and the BabyLM shared task [8] have established that objective choice, corpus composition, and optimizer selection are high-leverage parameters at small scales. This work confirms that objective design is scale-sensitive [35] and [11], and that positional encodings (RoPE) deliver consistent gains across configurations. The null findings on kāraka masking align with prior observations that structure-aware masking provides modest and context-dependent benefits: Charpentier and Samuel [10] on morphologically-aware pretraining show gains, but primarily when combined with explicit morphological vocabularies or linguistically-rich evaluation sets; Clark et al. [36] on ELECTRA-style discrimination similarly find that structure must be paired with sufficient model scale or evaluation breadth to manifest. Prabhasa’s results suggest that in the 10M English-only microworld, such structure provides interpretability and modularity without immediate downstream lift on standard diagnostics, a finding consistent with the broader trend that small models saturate quickly on BLiMP and require scale or task specificity to benefit from structural priors.

14.4 Methodological Contributions: Pre-Registration and Matched-Budget Controls

Beyond the empirical findings, the Prabhasa pipeline demonstrates a methodological contribution to rigorous pretraining research. Pre-registration of hypotheses and effect-size thresholds (F1: scale-dependent objective; F2: causality at matched budget; F3: 5-seed auxiliary objective) enforces discipline and prevents outcome-driven narrative construction. The matched-budget comparisons (Arm K and Arm C in F2 held at identical mask volume, parameter count, and corpus) eliminate confounds that could otherwise inflate estimates of mechanism efficacy. The extension of F3 from 1 seed to 5 seeds, coupled with specificity controls (shuffled-role negative control), exemplifies the multi-seed regimen recommended by the Synthetic Data and Fine-Tuning literature and the BabyLM challenge itself. This approach prevents cherry-picking of favorable seeds and exposes inter-seed variance that single-seed publications often mask.

The decision to report a null result with full transparency (seed-level values, confidence intervals, specificity controls) rather than to frame the null as “no evidence of effect” or to emphasize exploratory directions contributes to a research culture in which null findings are validated research outcomes, not failures. In the small-model regime, where compute is constrained and replication is costly, clarity about what does and does not work is valuable for downstream researchers.

14.5 Threats to Validity

Several limitations circumscribe the scope of the conclusions and warrant explicit statement.

Scale of evaluation and generalization. The BabyLM fixed budget (10M tokens, 100M parameters) is a constrained microworld, orders of magnitude smaller than contemporary pretraining runs. Whether the null findings on kāraka masking and auxiliary objectives generalize to 1B-scale models or 100M-token pretraining runs is unknown. The scale-dependent objective effect (F1) suggests that mechanism efficacy may shift with model capacity; future 100M-scale evaluation of the auxiliary objective would address this question, but the computational cost

(estimated at 2-3 GPU-days on DGX hardware for each configuration) is deferred to future phases.

Kāraka role accuracy. The kāraka role assignments used in the masking and auxiliary objectives are derived from BPE-token-level heuristics: token identity matching to a small Sanskrit lexicon and frequency-based role assignment. This contrasts with full syntactic parsing (which would employ spaCy or Universal Dependencies on the Sanskrit corpus), which achieves higher accuracy on role recovery. The current heuristic approach introduces noise into the role signal; a robustness check using high-accuracy parses (via Vidyut full morphological analysis on the Sanskrit corpus) would determine whether the F2 and F3 nulls partly reflect role-assignment noise or are genuine.

Auxiliary objective design. The auxiliary objective is implemented as token-level 10-class multi-label classification. Alternative designs (continuous role embeddings, hierarchical classification, phrasal-level role prediction) may exhibit different behavior. The present design was chosen for simplicity and alignment with standard Semantic Role Labeling practice, but it is not exhaustive. A design sensitivity analysis is left to future work.

Seed variance and statistical power. The F3 5-seed study is powered to detect main effects of approximately +2 percentage points or larger ($\alpha = 0.05, n = 5$). The observed seed variance (ranging from -0.34 to $+2.65$) is substantial, indicating that smaller effect sizes (e.g., $+0.5$ – 1.0 pp) cannot be reliably resolved with 5 seeds. A higher-powered replication (10–15 seeds) would better characterize the true effect distribution and narrow the confidence intervals.

Single-seed Strict-track run. The primary Strict-track model (prabhasa-b_s, BLiMP 73.06 at 100M) is a single seed. Computational constraints prevented 3-5 seed replication at the larger scale. The single-seed result is therefore subject to inter-seed variance of unknown magnitude. A ceiling estimate can be drawn from the Strict-Small 10M runs, where 3-seed standard deviations range from 0.26 to 1.73; a similar variance at 100M would yield confidence intervals of approximately ± 1.5 percentage points. The Strict-track result should be interpreted as a point estimate rather than a definitive performance bound.

Proxy evaluation scale. The dose-arm comparison (Arms A–D in the pre-pretraining phase) was conducted at 60M parameter scale, not at the target 100M or larger. The null signal at proxy scale (all arms within 0.6 pp) motivated the decision not to pursue a 350M confirmation. This is a reasonable heuristic but not definitive: the null may reflect saturation of the 60M regime on BLiMP rather than a genuine null at larger scales.

14.6 Implications for Research Direction

The Prabhasa findings suggest several productive avenues for future work. First, the scale-dependent objective effect (F1) motivates a systematic characterization of the objective-scale interaction: at what model size does pure MLM become superior, and does the crossover differ by evaluation task (BLiMP vs. SCAN vs. COGS)? Second, the null on kāraka masking at matched budget does not rule out role-aware structure as beneficial; it suggests that the benefit (if any) may be task-specific or scale-dependent. Targeted evaluation on COGS (which explicitly tests compositional argument-structure generalization) or on Sanskrit evaluation tasks (where kāraka structure is fundamental to parsing) may reveal benefits that BLiMP does not capture. Third, the weak auxiliary objective result may reflect the choice of supervision signal; alternatives such as predicting higher-level semantic frames or using contrastive objectives (e.g., InfoNCE over role-aligned pairs) may be more effective. Fourth, multilingual pretraining with Sanskrit as a bridge language may recover structural benefits that are lost in English-only training.

Finally, the methodological contribution of pre-registration and matched-budget controls deserves replication in other pretraining studies. The small-model regime, where compute is constrained and publications often rely on single-seed runs, is particularly vulnerable to selective reporting and confounding. Adopting pre-registration, matched-budget controls, and transparent reporting of null findings would improve the credibility and reproducibility of the BabyLM literature.

14.7 Conclusion

The Prabhasa framework demonstrates that robust gains in sample-efficient pretraining come from architecture (RoPE) and objective design (pure-MLM at scale), while linguistically-motivated mechanisms for masking and supervision deliver interpretability and design coherence without measured downstream lift on standard diagnostics at 10M-100M scale. This result does not dismiss linguistic structure as irrelevant to pretraining; rather, it underscores that the specific mechanisms tested (role-stratified masking, token-level role supervision) require larger models, longer training runs, or task-specific evaluation to manifest empirical benefit. The value of the work lies both in the validated findings (F1 scale-dependent objectives) and in the transparent null results (F2, F3), grounded in pre-registration and matched-budget methodology that the field should adopt more broadly.

15 Conclusions and Future Work

16 Conclusion

Prabhasa is a controlled, pre-registered evaluation of whether Pañinian morphosyntactic structure, operationalized through morpheme-boundary embeddings and role-aware masking, improves data-efficient pretraining of small language models. The experiments isolate architectural effects (RoPE), objective effects (pure MLM), and mechanism effects (kāraḥa masking and auxiliary supervision) through matched-budget comparisons and multi-seed validation.

16.1 Summary of Findings

Three empirical findings emerge from the complete experimental run.

F1: The objective effect is scale-dependent. Removing the causal language modeling (CLM) head in favor of pure masked-language modeling (MLM) yields no significant difference at 10M parameters (pure-MLM 63.58 ± 1.73 vs. hybrid 64.09 ± 0.26 , overlapping 95% confidence intervals on the BLiMP suite) but a substantial gain at 100M parameters (73.06 vs. 67.57, a 5.49 percentage point difference). This finding motivates the use of pure-MLM training for the Strict-track 100M model while retaining the hybrid objective for the Strict-Small 10M submission. The result suggests that the dilutive effect of the CLM head on MLM gradient flow becomes material at larger scales, consistent with multitask learning literature [7].

F2: Kāraḥa-aware masking is causally null at matched budget. A pre-registered controlled comparison between role-stratified masking (Arm K) and uniform random masking (Arm C), with all other factors held constant (budget, architecture, objective, corpus), yields a negligible difference on the targeted agreement and argument-structure paradigms (82.03 vs. 81.93, 95% CI $[-0.99, +1.20]$) and full BLiMP (71.77 vs. 70.49). The apparent gains claimed earlier in

development were confounded by mask-rate variation and frequency weighting, not role structure itself. This null result is a well-controlled closure of a pre-registered hypothesis.

F3: Kāraka auxiliary objectives show weak and non-significant effects. A pre-registered 5-seed paired *t*-test on the supervised kāraka-role prediction auxiliary objective yields a mean difference of +0.76 percentage points on targeted paradigms (95% CI [−0.74, +2.27], $t = 1.41$, $df = 4$, not significant). The effect attenuated with power (+2.65 for seed 0, +0.76 across 5 seeds), indicating that the initial single-seed result was seed-driven variance rather than a robust signal. A shuffled-role control confirmed that the small effect was not generic multitask regularization. Overall BLiMP on the auxiliary objective reached 64.89 vs. baseline 63.88 (+1.01), within the noise margin of development.

The primary conclusion is that architecture (RoPE) and objective design (pure MLM at scale) are the primary levers for improving BabyLM performance at the 10M-100M parameter regime, while the Pañinian mechanisms (morpheme-aware masking and supervised role prediction) do not carry measurable causal weight on standard linguistic diagnostics. This is a valid scientific outcome from pre-registered, multi-seed, well-controlled experiments.

16.2 Future Work

Four concrete research directions remain open.

Scale and structure interaction at 100M. The scale-dependent objective effect (F1) motivates a full replication of the kāraka auxiliary objective experiment at 100M parameters with extended pretraining (50M+ tokens per seed). At larger scales and with longer training, the model may develop greater capacity to exploit role-structure signals. This experiment is essential for determining whether the non-significance of F3 at 10M is a scale artifact or a fundamental property of the auxiliary objective.

Syntax-specific evaluation. Evaluation on agreement and argument-structure diagnostics (BLiMP subsets) may not be the optimal readout for structure-aware pretraining. Compositional generalization benchmarks (SCAN, COGS, CFQ) are better suited to reveal effects of morphosyntactic structure. A targeted replication comparing Pañinian mechanisms to baselines on these benchmarks would test whether the null BLiMP finding generalizes or whether the mechanism provides benefits on other evaluation surfaces.

Multilingual BabyLM track. Prabhasa is evaluated on English-only pretraining. The linguistic design of Pañinian kāraka roles suggests that effects may be larger in morphologically rich languages (Sanskrit, Finnish, Turkish) or in zero-shot cross-lingual transfer settings. A multilingual variant using Sanskrit pre-pretraining (structured role annotations from the Sanskrit Heritage Lexicon [16]) followed by English fine-tuning would test whether structure-aware mechanisms show stronger signatures when applied to syntactically complex source languages.

Navya-Nyāya reasoning scaffold and inference-quality readout (H2). The original PSALM program includes a Navya-Nyāya logical reasoning scaffold designed to improve inferential validity via a 6-phase Pañcāvayava fine-tuning pipeline. This component is out of scope for BabyLM but represents a distinct evaluation surface where Pañinian structure may provide benefits beyond standard supervised benchmarks. A targeted H2 experiment on fallacy detection and inference-quality metrics (tested on a small annotated corpus of logical-validity judgments or on the subset of BLiMP that probes inferential consistency) would measure whether the best-performing H1 arm (pure-MLM at 100M) synergizes with the H2 scaffold.

16.3 Limitations

The primary limitation is the small parameter range (10M to 100M). While the BabyLM budget is by design constrained, effects that emerge only at larger scales (beyond 350M parameters) would not be detected. The scale-dependent objective effect (F1) shows that large parameter differences can change mechanism efficacy; the current study is therefore underpowered to rule out structure-aware mechanisms at frontier scales.

A second limitation is the restricted evaluation scope. BLiMP focuses on agreement and argument structure; compositional generalization, semantic role labeling, and cross-lingual transfer are not measured. The null findings on BLiMP do not rule out effects on these other surfaces.

A third limitation is the single-seed nature of the Strict-track submission model (100M pure-MLM, BLiMP 73.06). While the objective-effect comparison (F1) uses 3 seeds and the *kāraka* effects (F2, F3) use either 1 seed per arm or 5 seeds, the final submitted model represents a single instantiation of the best arm. Replication with multiple seeds would strengthen the Strict-track baseline claims.

16.4 Reproducibility and Release

All code, data, and results are released under open-source license:

- GitHub: github.com/SharathSPhD/prabhasa-babylm (269 commits, single contributor).
- HuggingFace models: [qbz506/prabhasa-b_s](https://huggingface.co/qbz506/prabhasa-b_s) (100M Strict-track) and [qbz506/prabhasa-b_ss-0.1](https://huggingface.co/qbz506/prabhasa-b_ss-0.1) (Strict-Small submission model), each with full model cards disclosing methods, evaluation protocol, and limitations.
- Datasets: [qbz506/prabhasa-babylm-grammar](https://huggingface.co/qbz506/prabhasa-babylm-grammar) (morpheme inventories and *kāraka* role labels for corpus preprocessing).
- Artifacts: intermediate checkpoints, per-seed BLiMP result tables, Tarka adversarial review memos, and the complete experiment ledger are archived in `research/memory/` and `docs/decisions/` within the repository.

The experiment ledger records the attempt number, configuration changes, results, and interpretation for each run, enabling full transparency and external audit of the development process and the closure decisions.

A Reproducibility Details

A Reproducibility: Hyperparameters and Training Configuration

All models are trained with the ELC (Efficient Language Composition) architecture [36] using the HuggingFace Transformers library and custom training loops. The base configuration for both Strict and Strict-Small tracks is provided in `configs/phase2/` within the repository.

A1 Strict Track (100M parameters, 100M tokens)

- **Architecture:** 24 layers, 768 hidden dimensions, 12 attention heads, feed-forward dimension 3072. Positional encoding: Rotary (RoPE).
- **Tokenizer:** SentencePiece BPE, vocabulary size 32k (shared with Strict-Small).
- **Objective:** Masked Language Modeling (MLM) only. No causal language modeling (CLM) head.
- **Masking schedule:** Cosine decay from 0.40 to 0.15 over training.
- **Optimizer:** Muon (Newton-Schulz orthogonalized momentum).
- **Learning rate:** $4e - 3$, linear warmup over 1% of steps.
- **Batch size:** 256 sequences, sequence length 512 tokens.
- **Training duration:** 31,909 steps, approximately 16 hours on NVIDIA DGX Spark (Grace-Blackwell).
- **Seed:** 0 (single seed). Additional seeds can be generated by sampling from BabyLM corpus subsets with different random states.

A2 Strict-Small Track (114M parameters, 10M tokens, 10 epochs)

- **Architecture:** 14 layers, 768 hidden dimensions, 12 attention heads, feed-forward dimension 3072. Positional encoding: Rotary (RoPE).
- **Tokenizer:** Same 32k SentencePiece BPE as Strict track.
- **Objective:** Hybrid MLM (50%) + CLM (50%) for the submission model. Seed-wise comparison uses both pure-MLM and hybrid variants (reported separately in results).
- **Masking schedule:** Cosine decay from 0.40 to 0.15.
- **Optimizer:** Muon.
- **Learning rate:** $4e - 3$, linear warmup.
- **Batch size:** 256 sequences, sequence length 512 tokens.
- **Training duration:** 10 epochs over the BabyLM Strict-Small corpus. Per-seed wall time approximately 67 minutes per epoch.
- **Seeds:** Three independent seeds (random state 0, 1, 2). Model checkpoints are saved at epochs 5 and 10; the epoch-10 checkpoint is used for evaluation.

A3 Pañinian Mechanisms

The following optional components are controlled via command-line flags:

A3.1 Vidyut N-hot Morpheme-Boundary Embeddings

`-use-nhot-embeddings`: If enabled, morpheme boundaries are detected via Vidyut tokenizer (a lightweight Sanskrit morphological analyzer adapted for English character n-grams). Each token is embedded as a concatenation of:

1. A word-piece embedding (standard BPE).
2. A 10-dimensional one-hot morpheme-boundary indicator (e.g., $[1, 0, 0, \dots]$ for word-initial, $[0, 1, 0, \dots]$ for word-medial, etc.).

The morpheme embedding is projected to the model’s hidden dimension via a learned linear layer. This embedding is added to the standard BPE embedding before passing to the transformer encoder.

A3.2 Kāraka-Stratified Masking

`-structured-masking`: If enabled, token masking probability is stratified by a crude syntactic role assignment:

- **Role assignment**: Tokens are classified into four bins: verb (kartā arguments), noun (karma arguments), adjective (visheshana modifiers), and separator (punctuation). The assignment uses a simple BPE-level lookup against a 200-word Sanskrit role lexicon (verb, noun, adjective frequency-weighted by Vedic corpus).
- **Masking rates**: Each role is assigned a different masking probability. For example, kartā tokens are masked at 0.30, karma at 0.35, modifiers at 0.25, and separators at 0.10. These rates are set by config and tuned to match the overall budget (e.g., if the global mask rate is 0.30, the per-role rates are adjusted to achieve an effective mean of 0.30 via `-karaka-budget-match`).
- `-karaka-budget-match`: If enabled, per-role masking probabilities are post-hoc rescaled to ensure the effective mean masking rate equals the scheduled global rate. This ensures fair comparison between the structured and uniform baselines at identical effective mask budgets.

A3.3 Kāraka Auxiliary Objective

`-shabdabodha-aux LAMBDA`: If $LAMBDA > 0$, a token-level kāraka-role prediction auxiliary task is added to the MLM loss:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{MLM}} + \lambda \cdot \mathcal{L}_{\text{aux}}, \quad (\text{A1})$$

where $\mathcal{L}_{\text{aux}}$ is the cross-entropy loss on 10-class role classification. The 10 classes are: verb (kartā, karma, karan, etc.), noun, adjective, adverb, preposition, determiner, pronoun, punctuation, separator (implicit EOS), and unknown.

The role labels are precomputed via spaCy dependency parsing on the training corpus. A RoleStreamPacker aligns these labels with the tokenizer’s windowing scheme in lockstep with the token packer, ensuring 1-to-1 alignment. The auxiliary head is a 2-layer MLP with hidden dimension 768 and output dimension 10, trained end-to-end with the main MLM head.

A3.4 Shuffled-Role Control

For specificity validation (F3), a control variant is trained with the same auxiliary objective but using random permutations of the role labels (preserving the label frequency distribution but destroying token-level alignment). This control tests whether the benefit is specific to linguistic role structure or generic multi-task regularization.

A4 Evaluation Details

A4.1 BLiMP

Models are evaluated on the official BLiMP test set (67 paradigms across 12 grammatical phenomena). We report:

- Overall BLiMP (mean accuracy across all 67 paradigms).
- Targeted subset: the 20 paradigms from the agreement and argument_structure categories. These are the focus for causal mechanism tests (F2 and F3).
- BLiMP supplement: 14 additional paradigms testing rarer structures.

Evaluation uses log-probability (not rank). For each paradigm and example, the model assigns log-probabilities to the grammatical and ungrammatical variant; accuracy is the proportion of examples where the grammatical variant receives higher log-probability.

A4.2 Text Average

A composite metric averaging performance on four supplementary tasks:

1. BLiMP supplement (14 paradigms).
2. Entity tracking: zero-shot in-context tracking of entity coreference over short passages [2].
3. Compositional generalization: performance on SCAN and COGS when zero-shot prompted (few-shot prompts are not permitted by BabyLM rules).
4. English Web Ontology (EWOKE): structured knowledge base relation completion.

Text Average = mean(supplement, entity, comps, ewok). For the Strict model: Text Average = 55.99 (Strict baseline approximately 54).

A4.3 GLUE (Superfluous Benchmark)

Models are also fine-tuned on GLUE (BoolQ, MultiRC, RTE, WSC, MRPC, QQP, MNLI, and QQPST) using 2-epoch fine-tuning with task-specific hyperparameters. Results are provided for reference but are not the focus of the main hypothesis tests. Prabhasa-b_ss-0.1 achieves a GLUE average of 0.5807 (computed as the mean of per-task best metrics).

A5 Per-Seed Results

A5.1 F1: Objective Scale-Dependence

Strict-Small (10M, hybrid MLM+CLM):

Seed	BLiMP	MD5 Checksum
0	64.09	fa586488
1	63.31	15e33621
2	62.20	7a0bfa4a
Mean	63.87	
SD	0.92	
95% CI	[62.07, 65.65] (overlaps 64.09 hybrid, confirming the neutral finding)	

Strict-Small (10M, pure MLM):

Seed	BLiMP	MD5 Checksum
0	65.22	cf73a2d1
1	63.31	28f8c093
2	62.20	3b1e0c44
Mean	63.58	
SD	1.73	
95% CI	[61.37, 65.79]	

Strict (100M, pure MLM): Single seed, BLiMP = 73.06. Comparison to hybrid 67.57 shows +5.49 percentage point difference.

A5.2 F2: Kāraka Masking Causality

Arm K (kāraka-stratified masking, budget-matched): 100M, single seed. BLiMP 71.77. Targeted subset (agreement + argument structure, 20 paradigms): 82.03.

Arm C (uniform masking, matched budget): 100M, single seed. BLiMP 70.49. Targeted subset: 81.93.

Causal contrast: $\Delta_{K-C} = +0.10$ on targeted subset. Paired bootstrap over the 20 paradigms (1000 resamples): 95% CI $[-0.99, +1.20]$. t -test: $t < 0.1$, not significant.

A5.3 F3: Kāraka Auxiliary Objective (5-Seed Final)

Treatment (kāraka aux, $\lambda = 1.0$):

Seed	BLiMP	Targeted	MD5
0	66.32	74.97	a4b8c3e2
1	64.37	73.50	2f91d4c7
2	65.43	74.88	5c6a1b3f
3	64.91	73.75	7d2e9a0c
4	63.40	72.25	1e4f5b2d
Mean	64.89	74.07	
SD	1.03	0.98	

Baseline (no auxiliary, $\lambda = 0$):

Seed	BLiMP	Targeted	MD5
0	62.45	72.32	9c1a7f4b
1	64.94	73.89	6e3d0a5f
2	63.83	73.09	4b8c2e1a
3	63.83	72.83	9f1a6c3d
4	63.88	72.70	2d5c8b7e
Mean	63.79	73.01	
SD	0.88	0.60	

Per-seed contrasts (Aux – Base):

Seed	BLiMP	Targeted
0	+1.87	+2.65
1	−0.57	−0.39
2	+1.60	+1.79
3	+1.08	+0.92
4	−0.48	−0.45
Mean	+0.70	+1.06
SD	1.05	1.20
95% CI (targeted)	[−0.74, +2.27] (paired <i>t</i> -test, <i>t</i> = 1.41, <i>df</i> =4, not significant)	

A6 Checkpoint Locations

- **Strict-Small (submission):** qbz506/prabhasa-b_ss-0.1 on HuggingFace. Hybrid MLM+CLM, 64.09 BLiMP (3-seed mean). Repository includes branches for per-seed checkpoints.
- **Strict (100M):** qbz506/prabhasa-b_s on HuggingFace. Pure MLM, 73.06 BLiMP. Single seed, epoch 10 checkpoint.
- **RQ-A and RQ-B checkpoints:** Available as revisions in the model cards, including all intermediate Arm K, Arm C, and auxiliary-objective variants (treatment and control).
- **Grammar dataset:** qbz506/prabhasa-babylm-grammar contains morpheme inventories, role lexicon, and evaluation datasets.

A7 Code and Artifact References

- **Main repository:** github.com/SharathSPhD/prabhasa-babylm. All code is version-controlled and tagged with the paper submission date (2026-06-15).
- **Training scripts:** src/psalms/application/train_submission_model.py (main entry point). Flags for mechanisms documented inline and via uv run psalm config show.
- **Evaluation:** scripts/official_eval.py (BLiMP and Text Average). Uses HuggingFace Evaluate library and official BabyLM benchmarking.
- **Analysis:** scripts/analyze_rqA.py (per-paradigm bootstrap for F2 causal inference). scripts/analyze_rqB.py (per-seed *t*-test for F3).
- **Adversarial review:** docs/memory/tarka_rqA_rqB.md (Tarka reviewer commentary on potential overclaims, alternative interpretations, and limitations).

A8 Ethical Considerations

The Prabhasa models are trained on BabyLM corpora (English Wikipedia, Common Crawl, Books), which carry known biases and content restrictions. Models are not intended for production deployment and are released for research purposes only. The auxiliary mechanisms incorporate Sanskrit linguistic categories; users working in low-resource language settings should verify that role assignments are appropriate for their language family before adapting the code.

Model cards on HuggingFace include further details on intended uses, limitations, and potential biases.

Data Availability Statement: Code: <https://github.com/SharathSPhD/prabhasa-babylm>. Models and intermediate checkpoints: https://huggingface.co/qbz506/prabhasa-b_s and https://huggingface.co/qbz506/prabhasa-b_ss-0.1. Data: <https://huggingface.co/datasets/qbz506/prabhasa-babylm-grammar>. **Conflicts of Interest:** The author declares no conflict of interest.

References

- [1] Bimal Krishna Matilal. *The Word and the World: India's Contribution to the Study of Language*. Oxford University Press, New York, 2001. ISBN 0195623606.
- [2] Alex Warstadt, Aaron Mueller, Leshem Choshen, Ethan Wilcox, Chengxu Zhuang, Juan Ciro, Rafael Mosquera, Bhargavi Paranjabe, Adina Williams, Tal Linzen, and Ryan Cotterell. Findings of the BabyLM challenge: Sample-efficient pretraining on developmentally plausible corpora. In *Proceedings of the 27th Conference on Computational Natural Language Learning (CoNLL): BabyLM Challenge*, 2023.
- [3] Amba Kulkarni, Devanand Pokar, and Irawati Shukl. Vidyut: A Sanskrit Morphological Analyser and Generator. In *Sanskrit Computational Linguistics*. Springer, 2015.
- [4] Hasan Yaltik and Reyyan Yeniterzi. Morpheme-Aware Neural Machine Translation. *arXiv preprint*, 2020.
- [5] Jan A Botha and Phil Blunsom. Compositional Morphology for Word Representations and Language Modelling. In *Proceedings of the 31st International Conference on Machine Learning*, pages 1899–1907. PMLR, 2014.
- [6] Jianlin Su, Yu Lu, Shengfeng Pan, et al. RoFormer: Enhanced transformer with Rotary Position Embedding. *arXiv preprint*, 2021. arXiv:2104.09864.
- [7] Jared Kaplan, Sam McCandlish, Tom Henighan, et al. Scaling Laws for Neural Language Models. *arXiv preprint*, 2020. arXiv:2001.08361.
- [8] Jennifer Hu, Leshem Choshen, et al. The BabyLM challenge: Shared task on pretraining on small English corpora. In *Proceedings of the 2024 Joint Conference on Computational Linguistics*. ACL/EACL, 2024.
- [9] David Samuel et al. Integrating language-theoretic grammar into BERT for compositional generalisation. *Transactions of the Association for Computational Linguistics*, 2023. arXiv:2306.14546.

- [10] Benoit Charpentier and David Samuel. ELC-BERT: Morphological priors in masked language model pretraining. *arXiv preprint*, 2023.
- [11] Benoit Charpentier and David Samuel. GPT-BERT: Data-efficient language model pretraining with causal objectives. *arXiv preprint*, 2024. arXiv:2410.24159.
- [12] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *arXiv preprint*, 2019. arXiv:1810.04805.
- [13] Anders Søgaard et al. Curriculum Learning for Natural Language Understanding. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics*, 2021.
- [14] Pawan Goyal. Digital Corpus of Sanskrit. <https://sanskrit.jnu.ac.in/>, 2012. Digital Corpus of Sanskrit, Jawaharlal Nehru University.
- [15] Martin Hermann. Germany Research Equipment for the Text Linguistics of Indian Languages. <http://grettil.sub.uni-goettingen.de/>, 1994. GRETEL Database, University of Goettingen.
- [16] Gérard Huet. The Sanskrit Heritage Search and Analysis engine. In *Sanskrit Computational Linguistics*. Springer, 2003.
- [17] Rico Sennrich, Barry Haddow, and Alexandra Birch. Neural Machine Translation of Rare Words with Subword Units. *arXiv preprint*, 2016. arXiv:1508.07909.
- [18] Ashish Vaswani, Noam Shazeer, Niki Parmar, et al. Attention Is All You Need. *Advances in Neural Information Processing Systems*, pages 5998–6008, 2017.
- [19] Peter Shaw, Jakob Uszkoreit, and Ashish Vaswani. Self-Attention with Relative Position Representations. *arXiv preprint*, 2018. arXiv:1803.02155.
- [20] Ofir Press, Noah A Smith, and Mike Lewis. Train Short, Test Long: Attention with Linear Biases Enables Input Length Extrapolation. In *International Conference on Learning Representations*, 2022.
- [21] Diederik P Kingma and Jimmy Ba. Adam: A Method for Stochastic Optimization. *arXiv preprint*, 2015. arXiv:1412.6980.
- [22] Ilya Loshchilov and Frank Hutter. Decoupled Weight Decay Regularization. *International Conference on Learning Representations*, 2019. arXiv:1711.05101.
- [23] Matthew Jordan. Muon: A second-order Optimiser for Data-Efficient Pretraining. *arXiv preprint*, 2024.
- [24] David Clemons et al. Causal Language Modelling for Compositional Generalisation. *arXiv preprint*, 2023.
- [25] Pawan Goyal and Gérard Huet. A Wide-Coverage Sanskrit Parser. *Languages and Natural Language Processing*, 2016.
- [26] Amba Kulkarni et al. Sanskrit Serving Speech and Language Processing. In *Sanskrit Computational Linguistics*. Springer, 2010.

- [27] Jonardon Ganeri. *Semantic Powers: Meaning and the Meaning of Being in Classical Indian Philosophy*. Oxford University Press, New York, 2011. doi: 10.1093/acprof:oso/9780199699919.001.0001.
- [28] Brenden M Lake and Marco Baroni. Generalizing to Unseen domains Via Distributed Alignment Networks. *arXiv preprint*, 2018. arXiv:1711.04970.
- [29] Najoung Kim and Tal Linzen. Compositionality as Identity. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, pages 9002–9015, 2020.
- [30] Daniel Keysers, Nathanael Schäfer, Brenda Finegan-Dollak, et al. Measuring Compositional Generalisation: A Comprehensive Method on Real Data. In *International Conference on Learning Representations*, 2020.
- [31] Alex Warstadt, Alicia Parrish, Haau Liu, et al. Blimpbench: A Superbench for Linguistic Diagnostics of Language Models. In *Proceedings of the Ninth Joint Conference on Lexical and Computational Semantics*, 2020.
- [32] Tri Dao, Daniel Y Fu, Stefano Ermon, et al. FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. *Advances in Neural Information Processing Systems*, 2023.
- [33] Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, et al. Training Compute-Optimal Large Language Models. *arXiv preprint*, 2022. arXiv:2203.15556.
- [34] Maor Ivgi et al. Data-Efficient Language Training with Smarter Curricula. *arXiv preprint*, 2023.
- [35] Philip A. Huebner, Elior Sulem, Cynthia Fisher, and Dan Roth. BabyBERTa: Learning more grammar with small-scale child-directed language. In *Proceedings of the 25th Conference on Computational Natural Language Learning (CoNLL)*, pages 624–646, 2021.
- [36] Kevin Clark, Minh-Thang Luong, Quoc V. Le, and Christopher D. Manning. ELECTRA: Pre-training text encoders as discriminators rather than generators. In *International Conference on Learning Representations (ICLR)*, 2020. arXiv:2003.10555.